FinBen: An Holistic Financial Benchmark for Large Language Models

Qianqian Xie[a], **Weiguang Han**[b], **Zhengyu Chen**[b], **Ruoyu Xiang**[a], **Xiao Zhang**[a], **Yueru He**[a], **Mengxi Xiao**[b], **Dong Li**[b], **Yongfu Dai**[g], **Duanyu Feng**[g], **Yijing Xu**[a], **Haoqiang Kang**[e], **Ziyan Kuang**[l], **Chenhan Yuan**[c], **Kailai Yang**[c], **Zheheng Luo**[c], **Tianlin Zhang**[c], **Zhiwei Liu**[c], **Guojun Xiong**[j], **Zhiyang Deng**[i], **Yuechen Jiang**[i], **Zhiyuan Yao**[i], **Haohang Li**[i], **Yangyang Yu**[i], **Gang Hu**[h], **Jiajia Huang**[k], **Xiao-Yang Liu**[e], **Alejandro Lopez-Lira**[d], **Benyou Wang**[f], **Yanzhao Lai**[m], **Hao Wang**[g], **Min Peng**[b], **Sophia Ananiadou**[c], **Jimin Huang**[a]

[a]The Fin AI, [b]Wuhan University, [c]The University of Manchester, [d]University of Florida, [e]Columbia University, [f]The Chinese University of Hong Kong, Shenzhen, [g]Sichuan University, [h]Yunnan University, [i]Stevens Institute of Technology [j]Stony Brook University, [k]Nanjing Audit University, [l]Jiangxi Normal University, [m]Southwest Jiaotong University

Abstract

LLMs have transformed NLP and shown promise in various fields, yet their potential in finance is underexplored due to a lack of comprehensive evaluation benchmarks, the rapid development of LLMs, and the complexity of financial tasks. In this paper, we introduce FinBen, the first extensive open-source evaluation benchmark, including 36 datasets spanning 24 financial tasks, covering seven critical aspects: information extraction (IE), textual analysis, question answering (QA), text generation, risk management, forecasting, and decision-making. FinBen offers several key innovations: a broader range of tasks and datasets, the first evaluation of stock trading, novel agent and Retrieval-Augmented Generation (RAG) evaluation, and three novel open-source evaluation datasets for text summarization, question answering, and stock trading. Our evaluation of 15 representative LLMs, including GPT-4, ChatGPT, and the latest Gemini, reveals several key findings: While LLMs excel in IE and textual analysis, they struggle with advanced reasoning and complex tasks like text generation and forecasting. GPT-4 excels in IE and stock trading, while Gemini is better at text generation and forecasting. Instruction-tuned LLMs improve textual analysis but offer limited benefits for complex tasks such as QA. FinBen has been used to host the first financial LLMs shared task at the FinNLP-AgentScen workshop during IJCAI-2024, attracting 12 teams. Their novel solutions outperformed GPT-4, showcasing FinBen's potential to drive innovation in financial LLMs. All datasets, results, and codes are released for the research community[1].

1 Introduction

Recently, Large Language Models (LLMs) (Brown et al., 2020) such as ChatGPT[2] and GPT-4 (OpenAI, 2023a), have reshaped the field of natural language processing (NLP) and exhibited remarkable capabilities in specialized domains across mathematics, coding, medicine, law, and finance (Bubeck

---

[1] https://github.com/The-FinAI/PIXIU
[2] https://openai.com/chatgpt

et al., 2023). Within the financial domain, recent several studies (Xie et al., 2023a; Lopez-Lira and Tang, 2023; Li et al., 2023c; Xie et al., 2023b) have shown the great potential of advanced LLMs such as GPT-4 on financial text analysis and prediction tasks. While their potential is evident, a comprehensive understanding of their capabilities and limitations for finance remains largely unexplored. This is due to a lack of extensive evaluation studies and benchmarks, and the inherent complexities associated with the professional nature of financial tasks.

Existing financial domain evaluation benchmarks including PIXIU (Xie et al., 2023b), FinanceBench (Islam et al., 2023) and BizBench (Koncel-Kedziorski et al., 2023), have **limited Evaluation Tasks** and primarily **focus on Financial NLP Tasks** (As shown in Table 1). Most existing benchmarks cover only a small number of evaluation tasks and are centered on evaluating NLP capabilities, such as information extraction and QA. While PIXIU stands out by covering the highest number of tasks, it includes only one evaluation task in most categories. This narrow focus limits their ability to comprehensively evaluate LLMs across the diverse and complex landscape of financial applications, such as forecasting, risk management, and decision-making. It is insufficient for a thorough evaluation of LLM capabilities, especially in the financial area.

Table 1: Comparison of different financial benchmarks based on the number of tasks and datasets used, and the task number distribution across various aspects including information extraction (IE), textual analysis (TA), question answering (QA), text generation (TG), risk management (RM), forecasting (FO), and decision-making (DM).

| Benchmark	Language	Dataset	Task	IE	TA	QA	TG	RM	FO	DM
CFBenchmark	Chinese	8	7	1	3	✗	3	✗	✗	✗
Fin-Eva	Chinese	1	1	✗	✗	1	✗	✗	✗	✗
PIXIU	English	15	8	1	3	1	1	1	1	✗
FinanceBench	English	1	1	✗	✗	1	✗	✗	✗	✗
BizBench	English	8	5	2	✗	2	1	✗	✗	✗
FinBen	English	36	24	6	8	3	1	4	1	1

Figure 1: Evaluation datasets for FinBen, with colors indicating dataset sizes ranging from 100 to 4,000.

To bridge this gap, we propose FinBen, a novel comprehensive open-source evaluation benchmark developed through the collaborative efforts of experts in both computer science and finance. FinBen comprises 36 datasets spanning 24 financial tasks, meticulously organized to assess LLMs across seven critical aspects: information extraction (IE), textual analysis (TA), question answering (QA), text generation (TG), risk management (RM), forecasting (FO), and decision-making (DM). Each category targets specific skills of financial data handling and analysis, ensuring a thorough evaluation of LLMs and showcasing their proficiency in managing complex financial scenarios.

FinBen introduces several innovations compared to existing benchmarks: 1) **New Tasks**: FinBen introduces a significantly larger number of tasks and datasets, making it the most holistic benchmark for financial LLMs with the highest number of tasks and datasets. This extensive range provides a more robust evaluation of LLM capabilities in diverse financial contexts. 2) **Broader Coverage**: Covering seven aspects of the financial sector, FinBen is the first benchmark to include the evaluation of stock trading, which is the fundamental task in the financial sector, involving complex decision-making processes that impact market dynamics and investment strategies. 3) **New Evaluation Strategy**: FinBen is the first benchmark to include agent-based evaluation and Retrieval-Augmented Generation (RAG) based evaluation. These innovative strategies provide a more dynamic and realistic assessment of LLMs, reflecting their ability to interact with and retrieve relevant information from vast datasets. 4) **Novel Datasets**: FinBen proposes three novel open-source datasets of text summarization, QA, and stock trading tasks for the research community, pushing the boundaries of what LLMs can achieve and setting a new standard for dataset comprehensiveness. 5) **Empowering Financial LLMs Research**: Based on FinBen, we hosted the first shared task (see Appendix F for more details)

focused on financial LLMs at the FinNLP-AgentScen workshop during IJCAI-2024 [3]. This event attracted 12 teams, leveraging our benchmark to develop novel LLMs-based solutions within the financial domain. Remarkably, the proposed methods achieved superior performance compared to GPT-4, demonstrating the benchmark's potential to foster innovation and advance the state-of-the-art (SOTA) in financial LLMs.

Based on FinBen, we access 15 representative general LLMs such as GPT-4, ChatGPT, and the latest Gemini, and financial LLMs, and have the following findings: 1) **Superior Capabilities with Limitations**: While LLMs exhibit exceptional prowess in IE and textual analysis tasks, they underperform in areas necessitating advanced reasoning and complex IE, such as text generation and forecasting. 2) **Potential in Stock Trading**: SOTA LLMs have demonstrated considerable promise in stock trading applications. However, there remains significant room for improvement due to their limitations in reasoning and comprehensive forecasting abilities. 3) **Closed-Source Superiority**: Closed-source commercial LLMs continue to lead in performance within the financial domain. Specifically, GPT-4 excels in IE, text analysis, QA, and intricate stock trading tasks, while Gemini shows superior capabilities in text generation and forecasting. 4) **Open-Source Improvements and Limitations**: While open-source, instruction-tuned financial LLMs have shown notable enhancements in textual analysis and IE tasks, the advantages of instruction-tuning are less pronounced when it comes to complex tasks such as QA, text generation, and forecasting.

In summary, the main contributions of this paper are: 1) we propose FinBen, the first comprehensive open-sourced evaluation benchmark for LLMs in the financial domain, 2) we propose a novel taxonomy covering seven aspects for organizing financial evaluation tasks, 3) we propose three novel evaluation datasets for the research community, and 4) we conduct systematic evaluation of 15 LLMs using FinBen, showcasing their limitations and advantages, and highlighting directions for future work.

2 FinBen

In this section, we delve into the specifics of FinBen, detailing the evaluation taxonomy, data sources, and evaluation tasks.

2.1 The Taxonomy of Financial Evaluation Tasks

In the dynamic landscape of financial technology, evaluating the capabilities of LLMs necessitates a comprehensive and structured approach. We propose a novel taxonomy for financial evaluation tasks, categorizing and assessing LLMs across seven financial domains inspired by established taxonomies in financial tasks Cao (2022); Li et al. (2023b); Zhao et al. (2024): **Information Extraction (IE)**, **Textual Analysis (TA)**, **Question Answering (QA)**, **Text Generation (TG)**, **Risk Management (RM)**, **Forecasting (FO)**, and **Decision-Making (DM)**. **Information Extraction** focuses on identifying key entities and relationships within financial documents, transforming unstructured data into structured insights Costantino and Coletti (2008). **Textual Analysis** delves into content and sentiment analysis of financial texts, aiding in market trend understanding Loughran and McDonald (2020). **Question Answering** evaluates the model's ability to comprehend and respond to financial queries Maia et al. (2018). **Text Generation** assesses the production of coherent financial text La Quatra and Cagliero (2020). **Risk Management** involves evaluating creditworthiness, detecting fraud, and ensuring regulatory compliance Aziz and Dowling (2019). **Forecasting** predicts future financial trends, enabling strategic responses to market dynamics Abu-Mostafa and Atiya (1996). Finally, **Decision-Making** assesses the model's proficiency in making informed financial decisions, such as developing trading strategies and optimizing investment portfolios Paiva et al. (2019).

2.2 Data Sources

FinBen's evaluation tasks are drawn from three primary data sources: 1) open-sourced datasets from existing studies comprising open-sourced datasets originally released for non-LLM evaluation settings. Domain experts have designed diverse prompts and reformulated these datasets into instruction-response pairs, making them suitable for evaluating the zero-shot performance of LLMs. 2) datasets from existing evaluation benchmarks including datasets from existing evaluation benchmarks such as

---

[3]https://sites.google.com/nlg.csie.ntu.edu.tw/finnlp-agentscen

PIXIU. These datasets have already been transformed into the instruction tuning format, allowing for seamless integration and direct use in FinBen. 3) novel datasets introduced in this paper. These datasets are designed to address gaps in existing benchmarks and provide unique challenges for financial LLMs evaluation. Novel datasets include (As shown in Table 2): **EDTSum**: We develop the EDTSum dataset for the text summarization task, focusing on financial news. This dataset contains 2,000 news with their headlines as the ground-truth summarization manually selected and cleaned from Zhou et al. (2021). This dataset represents a significant contribution to evaluating the ability of LLMs to generate concise and informative summaries. **FinTrade**: We develop the FinTrade dataset specifically for the stock trading task, incorporating historical stock prices, news, and sentiment data. This dataset is the collection of 10 stocks over one year. Our contribution enables the assessments of LLMs in agent-based financial trading scenarios. **Regulations**: We develop the regulations dataset for the long-form QA task related to Over-the-Counter (OTC) derivatives and general financial matters within the European Union. It precisely maps these inquiries to relevant articles from the European Market Infrastructure Regulation (EMIR) and other directives, with updates ranging from 2013 to 2021. It tests LLMs' ability to interpret and apply complex regulatory frameworks in real-world scenarios.

2.3 Tasks

Table 2 and Figure 1 shows all tasks, datasets, data statistics, and evaluation metrics covered by FinBen[4].

Information extraction: It spans seven datasets across six information extraction tasks. *1) Named entity recognition* extracts entities like LOCATION, ORGANIZATION, and PERSON from financial agreements and SEC filings, using the NER (Alvarado et al., 2015) and FINER-ORD (Shah et al., 2023b) datasets. *2) Relation extraction* identifies relationships such as "product/material produced" and "manufacturer" in financial news and earnings transcripts with the FINRED dataset (Sharma et al., 2022). *3) Causal classification* discerns whether sentences from financial news and SEC filings convey causality using the SC dataset (Mariko et al., 2020). *4) Causal detection* identifies cause and effect spans in financial texts with the CD dataset (Mariko et al., 2020). *5) Numeric labeling* tags numeric spans in financial documents using the FNXL dataset Sharma et al. (2023), focusing on automating the assignment of labels from a large taxonomy to numeral spans in sentences. *6) Textual analogy parsing* involves identifying common attributes and comparative elements in textual analogies by extracting analogy frames, utilizing the FSRL dataset Lamm et al. (2018), which maps analogous facts to semantic role representations and identifies the analogical relations between them. The evaluation of these tasks is focused on the F1 score (Goutte and Gaussier, 2005), Entity F1 score (Derczynski, 2016), and the Exact Match Accuracy (EMAcc) metric (Kim et al., 2023).

Textual analysis: This encompasses eight classification tasks for evaluating LLMs. *1) Sentiment analysis* focuses on extracting sentiment information (positive, negative, or neutral) from financial texts, using three datasets: the Financial Phrase Bank (FPB) (Malo et al., 2014), FiQA-SA (Maia et al., 2018), and TSA (Cortis et al., 2017) dataset. *2) News headline classification* analyzes additional information, like price movements in financial texts, using the Headlines dataset (Sinha and Khandait, 2021). *3) Hawkish-Dovish classification* aims to classify sentences from monetary policy texts as 'hawkish' or 'dovish' focusing on the nuanced language and economic implications of financial texts, using the FOMC (Shah et al., 2023a) dataset. *4) Argument unit classification* categorizes sentences as claims or premises using the FinArg AUC dataset (Sy et al., 2023). *5) Argument relation detection* identifies relationships (attack, support, or irrelevant) between social media posts using the FinArg ARC dataset (Sy et al., 2023). *6) Multi-class classification* targets categorizing a variety of financial texts, including analyst reports, news articles, and investor comments, utilizing the MultiFin dataset (Jørgensen et al., 2023). *7) Deal completeness classification* predicts if mergers and acquisitions events are "completed" or remain "rumors" based on news and tweets, employing the MA dataset (Yang et al., 2020a). *8) ESG issue identification* focuses on detecting Environmental, Social, and Governance (ESG) concerns in financial documents using the MLESG dataset (Chen et al., 2023a). For all datasets, evaluation utilizes the accuracy and F1 Score.

---

[4]For detailed instructions of each dataset, please see Appendix C

Table 2: The tasks, datasets, data statistics, and evaluation metrics included in FinBen. We use only test data for evaluation. Datasets marked with an asterisk (*) are newly constructed by us, comprising 10.32% of the total data. EM Accuracy means the exact match accuracy.

Data	Task	Test	Evaluation	License
NER (Alvarado et al., 2015)	named entity recognition	980	Entity F1	CC BY-SA 3.0
FiNER-ORD (Shah et al., 2023b)	named entity recognition	1080	Entity F1	CC BY-NC 4.0
FinRED (Sharma et al., 2022)	relation extraction	1,068	F1, Entity F1	Public
SC (Mariko et al., 2020)	causal classification	8,630	F1,Entity F1	CC BY 4.0
CD (Mariko et al., 2020)	causal detection	226	F1,Entity F1	CC BY 4.0
FNXL (Sharma et al., 2023)	numeric labeling	318	F1,EM Accuracy	Public
FSRL (Lamm et al., 2018)	textual analogy parsing	97	F1, EM Accuracy	MIT License
FPB (Malo et al., 2014)	sentiment analysis	970	F1, Accuracy	CC BY-SA 3.0
FiQA-SA (Maia et al., 2018)	sentiment analysis	235	F1	Public
TSA (Cortis et al., 2017)	sentiment analysis	561	F1, Accuracy	CC BY-NC-SA 4.0
Headlines (Sinha and Khandait, 2021)	news headline classification	2,283	Avg F1	CC BY-SA 3.0
FOMC (Shah et al., 2023a)	hawkish-dovish classification	496	F1, Accuracy	CC BY-NC 4.0
FinArg-ACC (Sy et al., 2023)	argument unit classification	969	F1, Accuracy	CC BY-NC-SA 4.0
FinArg-ARC (Sy et al., 2023)	argument relation classification	496	F1, Accuracy	CC BY-NC-SA 4.0
MultiFin (Jørgensen et al., 2023)	multi-class classification	690	F1, Accuracy	Public
MA (Yang et al., 2020a)	deal completeness classification	500	F1, Accuracy	Public
MLESG (Chen et al., 2023a)	ESG Issue Identification	300	F1, Accuracy	CC BY-NC-ND
FinQA (Chen et al., 2021)	question answering	1,147	EM Accuracy	MIT License
TATQA (Zhu et al., 2021)	question answering	1,668	F1, EM Accuracy	MIT License
*Regulations	long-form question answering	254	ROUGE, BERTScore	Public
ConvFinQA (Chen et al., 2022b)	multi-turn question answering	1,490	EM Accuracy	MIT License
ECTSum (Mukherjee et al., 2022)	text summarization	495	ROUGE, BERTScore, BARTScore	Public
*EDTSum	text summarization	2000	ROUGE, BERTScore, BARTScore	Public
BigData22 (Soun et al., 2022)	stock movement prediction	1,470	Accuracy, MCC	Public
ACL18 (Xu and Cohen, 2018)	stock movement prediction	3,720	Accuracy, MCC	MIT License
CIKM18 (Wu et al., 2018)	stock movement prediction	1,140	Accuracy, MCC	Public
German (Hofmann, 1994)	credit scoring	1000	F1, MCC	CC BY 4.0
Australian (Quinlan, [n. d.])	credit scoring	690	F1, MCC	CC BY 4.0
LendingClub (Feng et al., 2023)	credit scoring	2,690	F1, MCC	CC0 1.0
ccf (Feng et al., 2023)	fraud detection	2,278	F1, MCC	(DbCL) v1.0
ccfraud (Feng et al., 2023)	fraud detection	2,097	F1, MCC	Public
polish (Feng et al., 2023)	financial distress identification	1,736	F1, MCC	CC BY 4.0
taiwan (Feng et al., 2023)	financial distress identification	1,364	F1, MCC	CC BY 4.0
ProtoSeguro (Feng et al., 2023)	claim analysis	2,381	F1, MCC	Public
travelinsurance (Feng et al., 2023)	claim analysis	3,800	F1, MCC	(ODbL) v1.0
*FinTrade	stock trading	3,384	CR, SR, DV, AV, MD	MIT License

Question answering. It includes 3 datasets from 2 QA tasks, challenging LLMs to respond to financial queries. *1) QA* focuses on solving questions through multi-step numerical reasoning with financial reports and tables, utilizing the FinQA (Chen et al., 2021) and TATQA (Zhu et al., 2021) dataset. *2) Multi-turn QA* is an extension of QA with multi-turn questions and answers based on financial earnings reports and tables, using the ConvFinQA dataset (Chen et al., 2022b). F1 score (Derczynski, 2016) and the Exact Match Accuracy (EMAcc) metric (Kim et al., 2023) are used to evaluate these tasks. *2) Long-form QA* involve presenting models with complex, detailed questions that require extensive and nuanced answers, often incorporating legal interpretations and practical applications. In our evaluation, we utilize our newly proposed Regulations dataset, which focuses on intricate questions and answers related to financial regulations like EMIR. We assess the model responses using ROUGE (Lin, 2004) and BERTScore (Zhang et al., 2019).

Text generation. This task assesses the models' ability to produce coherent and informative text. Our focus is on *text summarization*, utilizing the ECTSUM (Mukherjee et al., 2022) dataset for summarizing earnings call transcripts. We also propose a novel dataset, EDTSUM, specifically designed for condensing financial news articles into concise summaries, constructed from original data in (Zhou et al., 2021). Evaluation employs ROUGE (Lin, 2004), BERTScore (Zhang et al., 2019), and BART Score (Yuan et al., 2021) to measure alignment, factual consistency, and information retention between machine-generated and expert summaries.

Forecasting. The forecasting task challenges models to adaptively predict future market and investor behaviors from emerging patterns. We focus on the *1) Stock movement prediction* task, forecasting stock directions as either positive or negative, based on historical prices and tweets, utilizing three datasets: BigData22 (Soun et al., 2022), ACL18 (Xu and Cohen, 2018) and CIKM18 (Wu et al., 2018).

Risk management. It challenges LLMs to accurately identify, extract, and analyze relevant risk-related information, interpret numerical data, and understand complex relationships. We include 4 tasks: *1) Credit scoring* classifies individuals as "good" or "bad" credit risks using historical customer

data, employing datasets including: German (Hofmann, 1994), Australia (Quinlan, [n. d.]) and LendingClub (Feng et al., 2023). **2) Fraud detection** involve categorizes transactions as "fraudulent" or "non-fraudulent", using two datasets: ccf (Feng et al., 2023) and ccFraud (Feng et al., 2023). **3) Financial distress identification** aims to predict a company's bankruptcy risk, using the polish (Feng et al., 2023) and taiwan dataset (Feng et al., 2023). Note that the dataset name describes only the region of the company, and the content within the datasets is in English. **4) Claim analysis** anonymizes client data for privacy, labeling a "target" to indicate claim status, using two datasets: PortoSeguro (Feng et al., 2023) and travelinsurance (Feng et al., 2023). It is noticed that the dataset name such as German and taiwan, only indicates customer sources and all content is in English. F1 score and Matthews correlation coefficient (MCC) (Chicco and Jurman, 2020) are used for evaluating these tasks.

Decision-making. Strategic decision-making (Punt, 2017) evaluates the model's proficiency in synthesizing diverse information to formulate and implement trading strategies, a challenge even for experts. We innovatively introduce the SOTA financial LLM agent FinMem (Yu et al., 2023) to evaluate LLMs on the **stock trading** task. We construct the novel FinTrade dataset, containing 10 stocks, simulating real-world trading through historical prices, news, and sentiment analysis. Performance is measured by Cumulative Return (CR) (Ariel, 1987), Sharpe Ratio (SR) (Sharpe, 1998), Daily (DV) and Annualized volatility (AV) (Zhou et al., 2023), and Maximum Drawdown (MD) (Magdon-Ismail and Atiya, 2004), offering a comprehensive assessment of profitability, risk management, and decision-making prowess.

3 Evaluation

We evaluate the zero-shot (from our evaluation) and few-shots (results from previous papers) performance of 15 representative general LLMs and financial LLMs on the FinBen benchmark, including: 1) ChatGPT: A LLM developed by OpenAI. 2) GPT-4 (OpenAI, 2023b): The SOTA commercialized LLMs proposed by OpenAI. 3) Gemini Pro Team et al. (2023): A multimodal AI LLM with 50T parameters, released by Google. 4) LLaMA2-7/70B-chat Touvron et al. (2023b): An open-sourced instruction-following LLM with 7B and 70B parameters developed by MetaAI. 5) LLaMA3-8B[5]: An open-sourced LLMs developed by MetaAI, using more training data than LLaMA2. 6) ChatGLM3-6B Du et al. (2022): A conversational LLM with 6B parameters, jointly released by Zhipu AI and Tsinghua KEG. 7) Baichuan2-6B Baichuan (2023): An open-source LLM with 6B parameters, launched by Baichuan Intelligent Technology. 8) InternLM-7B Team (2023): An open-sourced 7B parameter base model tailored for practical scenarios, proposed by SenseTime. 9) Falcon-7B Almazrouei et al. (2023): A 7B parameter causal decoder-only LLM model trained on 1500B tokens of RefinedWeb enhanced with curated corpora. 10) Mixtral $8\times$7B Jiang et al. (2024): A LLM with the Sparse Mixture of Experts (SMoE) architecture. 11) Code LLaMA-7B Roziere et al. (2023): An open-source LLM model for generating programming code, launched by Meta AI with 7B parameters. 12) FinGPT Yang et al. (2023a): A 7B instruction finetuned financial LLM based on LLaMA 7B Touvron et al. (2023a) with sentiment analysis tasks. 13) FinMA-7B Xie et al. (2023b): A 7B instruction finetuned financial LLM based on LLaMA 7B with multiple NLP and forecasting tasks. 14) DISC-FinLLM Chen et al. (2023b): An open-sourced financial LLM, fine-tuned from Baichuan-13B-Chat Baichuan (2023). 15) CFGPT Li et al. (2023a): An open-source LLM, specifically designed for the financial sector and trained on Chinese financial datasets, which comprises 7B parameters.

Experimental Settings We set the maximum generation tokens for LLMs to 1024 and the batch size to 20,000 for all experiments. These experiments are exclusively conducted on 16 NVIDIA A100 80G GPUs, taking approximately 600 hours to complete. Including the GPT-4 API costs, the total expenditure amounts to approximately $51,000.

4 Results

Table 3 and Table 4 shows the performance of 12 representative LLMs on all datasets in the FinBen. We also report results of non-LLM methods (traditional methods) in Appendix G.

---

[5]https://llama.meta.com/llama3/

4.1 Information Extraction & Textual Analysis Results

As shown in Table 3, for IE tasks, GPT-4 demonstrates superior performance in named entity recognition tasks, including NER, FINER-ORD, and FinRED. InternLM 7B achieves the best results in causal classification (SC). However, for more complex information extraction tasks, such as causal detection (CD) and numerical understanding (FNXL and FSRL), even GPT-4's performance is limited, with Gemini showing only slightly better results, still falling short of expectations. Additionally, while financial domain-specific LLMs developed by instruction tuning such as FinMA 7B exhibit improvements over general domain LLMs such as LLaMA2 7B-chat, they continue to struggle with both named entity recognition and complex extraction tasks. These findings highlight significant opportunities for advancement in financial causal detection and numerical understanding for LLMs.

Table 3: The zero-shot and few-shot performance of different LLMs on the FinBen. All results via our evaluations are the average of three runs. "-" represents the result that is currently unable to yield due to model size or availability, and "*" represents the result from the previous paper.

Dataset	Metrics	Chat GPT	GPT 4	Gemini	LLaMA2 7B-chat	LLaMA2 70B	LLaMA3 8B	FinMA 7B	FinGPT 7b-lora	InternLM 7B	Falcon 7B	Mixtral 7B	CFGPT sft-7B-Full
NER	EntityF1	0.77*	**0.83***	0.61	0.18	0.04	0.08	0.69	0.00	0.00	0.00	0.24	0.00
FINER-ORD	EntityF1	0.28	**0.77**	0.14	0.02	0.07	0.00	0.00	0.00	0.00	0.00	0.05	0.00
FinRED	F1	0.00	**0.02**	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
SC	F1	0.80	0.81	0.74	0.85	0.61	0.69	0.19	0.00	**0.88**	0.67	0.83	0.15
CD	F1	0.00	0.01	**0.03**	0.00	0.01	0.00	0.00	0.00	0.00	0.00	0.00	0.00
FNXL	EntityF1	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
FSRL	EntityF1	0.00	0.01	**0.03**	0.00	0.01	0.00	0.00	0.00	0.00	0.00	0.00	0.00
FPB	F1	0.78*	0.78*	0.77	0.39	0.73	0.52	**0.88**	0.00	0.69	0.07	0.29	0.35*
	Acc	0.78*	0.76*	0.77	0.41	0.72	0.52	**0.88**	0.00	0.69	0.05	0.37	0.26*
FiQA-SA	F1	0.60	0.80	0.81	0.76	**0.83**	0.70	0.79	0.00	0.81	0.77	0.16	0.42*
TSA	RMSE↓	0.53	0.50	0.37	0.71	0.57	0.25	0.80	0.00	**0.29**	0.50	0.50	1.05
Headlines	AvgF1	0.77*	0.86*	0.78	0.72	0.63	0.60	**0.97**	0.60	0.60	0.45	0.60	0.61*
FOMC	F1	0.64	**0.71**	0.40	0.35	0.49	0.40	0.49	0.00	0.36	0.30	0.37	0.16*
	Acc	0.6	**0.69**	0.60	0.49	0.47	0.41	0.46	0.00	0.35	0.30	0.35	0.21*
FinArg-ACC	MicroF1	0.50	**0.60**	0.31	0.46	0.58	0.51	0.27	0.00	0.39	0.23	0.39	0.05
FinArg-ARC	MicroF1	0.39	0.40	**0.60**	0.27	0.36	0.28	0.08	0.00	0.33	0.32	0.57	0.05
MultiFin	MicroF1	0.59	**0.65**	0.62	0.20	0.63	0.39	0.14	0.00	0.34	0.09	0.37	0.05
MA	MicroF1	0.85	0.79	0.84	0.70	**0.86**	0.34	0.45	0.00	0.78	0.39	0.34	0.25
MLESG	MicroF1	0.25	**0.35**	0.34	0.03	0.31	0.12	0.00	0.00	0.14	0.06	0.17	0.01
FinQA	EmAcc	0.58*	**0.63***	0.00	0.00	0.06	0.00	0.04	0.00	0.00	0.00	0.00	0.00
TATQA	EmAcc	0.00*	0.13*	**0.18**	0.03	0.01	0.01	0.00	0.00	0.00	0.00	0.01	0.00
Regulations	Rouge-1	0.12	0.11	-	0.24	-	0.10	0.12	0.01	0.04	0.03	-	0.14
	BertScore	0.64	0.62	-	0.65	-	0.60	0.59	0.40	0.57	0.14	-	0.57
ConvFinQA	EmAcc	0.60*	**0.76***	0.43	0.00	0.25	0.00	0.20	0.00	0.00	0.00	0.31	0.01
EDTSUM	Rouge-1	0.17	0.20	**0.39**	0.17	0.25	0.14	0.13	0.00	0.13	0.15	0.12	0.01
	BertScore	0.66	0.67	**0.72**	0.62	0.68	0.60	0.38	0.52	0.48	0.57	0.61	0.51
	BartScore	-3.64	**-3.62**	-3.87	-3.99	-3.81	-4.94	-5.71	-7.23	-4.60	-6.1	-4.47	-7.08
ECTSUM	Rouge-1	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
	BertScore	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
	BartScore	-5.18	-5.18	-4.93	-5.18	**-4.86**	-5.18	-5.18	-5.18	-5.18	--5.18	-5.18	-5.18
BigData22	Acc	0.53	0.54	**0.55**	0.54	0.47	0.55	0.51	0.45	0.56	0.55	0.46	0.45
	MCC	-0.025	0.03	0.04	0.05	0.00	0.02	0.02	0.00	**0.08**	0.00	0.02	0.03
ACL18	Acc	0.50	**0.52**	0.52	0.51	0.51	0.52	0.51	0.49	0.51	0.51	0.49	0.48
	MCC	0.005	0.02	**0.04**	0.01	0.01	0.02	0.03	0.00	0.02	0.00	0.00	-0.03
CIKM18	Acc	0.55	**0.57**	0.54	0.55	0.49	0.57	0.50	0.42	0.57	0.47	0.42	0.41
	MCC	0.01	0.02	0.02	-0.03	-0.07	0.03	**0.08**	0.00	-0.03	-0.06	-0.05	-0.07
German	F1	0.20	0.55	0.52	**0.57**	0.17	0.56	0.17	0.52	0.41	0.23	0.53	0.53
	MCC	-0.10	-0.02	0.00	**0.03**	0.00	0.05	0.00	0.00	-0.30	-0.07	0.00	0.00
Australian	F1	0.41	**0.74**	0.26	0.26	0.41	0.26	0.41	0.38	0.34	0.26	0.26	0.29
	MCC	0.00	**0.47**	0.00	0.00	0.00	0.00	0.00	0.11	0.13	0.00	0.00	-0.10
LendingClub	F1	0.20	0.55	0.65	**0.72**	0.17	0.10	0.61	0.00	0.59	0.02	0.61	0.05
	MCC	-0.10	-0.02	**0.19**	0.00	0.00	-0.15	0.00	0.00	0.15	-0.01	0.08	0.01
ccf	F1	0.20	0.55	0.96	0.00	0.17	0.01	0.00	1.00	**1.00**	0.10	0.00	0.00
	MCC	-0.10	-0.02	-0.01	**0.00**	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
ccfraud	F1	0.20	0.55	**0.90**	0.25	0.17	0.36	0.01	0.00	0.57	0.62	0.48	0.03
	MCC	-0.10	-0.02	0.00	-0.16	0.00	-0.03	-0.06	0.00	-0.13	-0.02	**0.16**	0.01
polish	F1	0.20	0.55	0.86	0.92	0.17	0.83	0.92	0.30	0.92	0.76	**0.92**	0.40
	MCC	-0.10	-0.02	**0.14**	0.00	0.00	-0.06	-0.01	0.00	0.07	0.05	0.00	-0.02
taiwan	F1	0.20	0.55	**0.95**	0.95	0.17	0.26	0.95	0.60	0.95	0.00	0.95	0.70
	MCC	-0.10	-0.02	0.00	**-0.01**	0.00	-0.07	0.00	-0.02	-0.01	0.00	0.00	0.00
portoseguro	F1	0.20	0.55	0.95	0.01	0.17	0.94	0.04	**0.96**	0.96	0.95	0.72	0.00
	MCC	-0.10	-0.02	0.00	-0.05	0.00	-0.01	**0.01**	0.00	0.00	0.00	0.01	0.00
travelinsurance	F1	0.20	0.55	0.00	0.00	0.17	0.00	0.00	**0.98**	0.89	0.77	0.00	0.03
	MCC	-0.10	-0.02	0.00	0.00	0.00	0.00	0.00	0.00	**0.12**	-0.03	0.00	0.01

Regarding TA tasks, instruction fine-tuned models like FinMA 7B exhibit the best performance in sentiment analysis tasks, including FPB, FiQA-SA, and Headlines. However, the generalization ability of FinMA 7B is limited due to the diversity of TA tasks in the financial domain. It performs even worse than general domain LLMs such as LLaMA2-7B-chat on other TA tasks, where GPT-4, Gemini, and LLaMA2 70B show superior results. This underscores the limitations of instruction fine-tuned models, which may also be constrained by the parameter size of their backbone models.

Models tailored for the Chinese language, such as CFGPT sft-7B-Full, which is fine-tuned on Chinese financial data, exhibit limited improvement on some datasets and even a decline in performance on others like MultiFin compared to its base model InternLM 7B. This trend suggests a language-based discrepancy, indicating that fine-tuning with Chinese data may adversely affect performance on English tasks. These findings underscore the complexities of cross-lingual adaptation in model training, highlighting the challenges in achieving consistent performance across different languages.

4.2 Question Answering & Text Generation Results

In the QA tasks, closed-source commercial LLMs like GPT-4 and Gemini continue to lead across all datasets. While FinMA 7B shows improvement over its backbone models, it remains limited by model size and exhibits bottlenecks in numeric reasoning ability. For the regulations dataset, which is the first intersection dataset requiring both financial and legal knowledge, GPT-4 demonstrates its broad knowledge coverage effectively.

In the TG tasks, Gemini emerges as the frontrunner on the EDTSUM abstractive text summarization dataset, illustrating its prowess in generating coherent summaries. Nevertheless, all models face challenges with extractive summarization, which demands the generation of precise label sequences for sentences. Among open-source LLMs, LLaMA2 70B stands out in text summarization. Conversely, CFGPT sft-7B-Full consistently shows a decrease in performance compared to its foundational model, InternLM 7B.

4.3 Forecasting & Risk Management Results

For forecasting, it is crucial to acknowledge that all LLMs fail to meet expected outcomes and lag behind traditional methodologies. This consistent observation with existing studies Xie et al. (2023b) underlines a notable deficiency in LLMs' capacity to tackle forecasting as effectively as traditional methods. Even the best-performing models, such as GPT-4 and Gemini, only perform slightly better than random guessing. This reveals significant potential for enhancement in LLMs, including industry leaders like GPT-4 and Gemini, particularly in forecasting tasks that demand complex reasoning abilities.

In RM tasks, such as credit scoring, fraud detection, and identifying financial distress, data often exhibit significant imbalances. Instances representing individuals with low credit scores, those prone to fraud, and companies at risk of financial distress constitute only a small percentage of the overall dataset. In such scenarios, LLMs with low instruction-following abilities (such as LLaMA2-7B-chat and LLaMA2-70B) tend to classify all cases into a single class, resulting in an MCC score of 0. These tasks, with tabular inputs and highly imbalanced distribution, pose a significant challenge for LLMs in the financial domain.

4.4 Decision Making Results

The comparative analysis of various LLMs on the complex task of stock trading, is presented in Table 4[6]. This task requires models to understand, summarize, and reason with multimodal financial data (texts and time series), leading to sophisticated trading decisions that necessitate a range of skills, from fundamental comprehension and summarization to reasoning and decision-making.

Among the evaluated LLMs, GPT-4 distinguishes itself by achieving the highest Sharpe Ratio (SR) over 1, indicating superior investment performance through optimal risk-return balance. It also records the minimal Maximum Drawdown (MDD), suggesting effective limitation of potential losses, thereby offering a more secure investment avenue compared to other models, including those using reinforcement learning methods like DQN, PPO, and A2C, which show significantly lower SR and higher MDD.

Tables 4 and 9 reinforce these findings, highlighting GPT-4's exceptional performance in this challenging domain. Additional results and analyses from these models in Table 5 contrast their performances with the traditional *Buy & Hold* strategy, which considerably lags behind.

In contrast, ChatGPT exhibits significantly lower performance metrics, indicating limitations in its financial decision-making capabilities. Gemini, on the other hand, secures the position of second-best

---

[6]For detailed trading performance, please see Appendix E

Table 4: The average trading performance (95% Confidence Interval) comparison for different LLMs across 10 stocks. The results include large LLMs only ($\geq 70B$), as models with smaller contexts have difficulty understanding the instructions and producing a static strategy of holding.

Model	CR (%)↑	SR↑	DV (%)↓	AV (%)↓	MD (%)↓
Buy & Hold	-4.00 ± 22.39	0.02 ± 0.87	3.59 ± 1.34	56.43 ± 21.00	30.67 ± 17.48
GPT-4	**28.19 ± 25.27**	**1.51 ± 1.08**	2.52 ± 1.30	39.88 ± 20.66	**18.34 ± 9.77**
GPT-4o	-5.54 ± 19.12	-0.19 ± 0.84	2.73 ± 1.30	43.62 ± 20.67	29.96 ± 18.89
GPT3.5-Turbo	4.48 ± 22.23	0.15 ± 0.82	2.84 ± 1.47	45.39 ± 23.35	28.83 ± 15.40
llama2-70B	4.02 ± 24.65	0.52 ± 1.48	2.18 ± 1.28	34.86 ± 20.38	25.55 ± 16.83
llama3-70B	-2.57 ± 22.63	-0.04 ± 1.19	2.71 ± 1.54	43.42 ± 24.65	29.31 ± 15.57
gemini	14.95 ± 28.04	1.03 ± 1.24	**2.17 ± 1.39**	**34.67 ± 22.23**	20.13 ± 11.36

Table 5: Traditional model performances on stock trading.

Model	Cumulative Return	Sharpe Ratio	Standard Deviation	Annualized Volatility	Max Drawdown
A2C	-4.2232	-0.2586	2.7522	43.6898	30.5819
PPO	-0.5586	0.0085	2.7531	43.7048	28.9496
DQN	-2.9924	-0.1656	2.7486	43.6319	31.78

performer, showcasing lower risk and volatility compared to GPT-4, yet maintaining commendable returns. When considering open-source models, LLaMA-70B, despite its lower volatility, yields the least profit among the LLMs, highlighting a trade-off between risk management and profitability.

For smaller models with parameters less than 70 billion, a marked inability to adhere to trading instructions consistently across transactions is noted. This is attributed to their limited comprehension, extraction capabilities, and constrained context windows. This limitation underscores the critical challenges smaller LLMs face in tasks requiring intricate financial reasoning and decision-making, thereby spotlighting the necessity for more advanced models to tackle decision making tasks effectively.

5 Conclusion

In this work, we present FinBen, a comprehensive benchmark specifically designed to evaluate LLMs in the financial domain. FinBen includes 36 diverse datasets spanning 24 tasks, meticulously organized to assess LLMs across seven critical aspects: information extraction, textual analysis, question answering, text generation, risk management, forecasting, and decision-making. This breadth of coverage sets FinBen apart from existing financial benchmarks, enabling a more robust and nuanced evaluation of LLM capabilities. Our evaluation of 15 LLMs, including GPT-4, ChatGPT, and Gemini, reveals their key advantages and limitations, highlighting directions for future work. Looking ahead, FinBen aims to continuously evolve, incorporating additional languages and expanding the range of financial tasks to further enhance its applicability and impact.

Limitations: We acknowledge several limitations that could impact FinBen's effectiveness and applicability. The restricted size of available datasets may affect the models' financial understanding and generalization across various contexts. Computational constraints limited our evaluation to the LLaMA 70B model, potentially overlooking the capabilities of larger models. Additionally, the tasks are based on American market data and English texts, which may limit the benchmark's applicability to global financial markets. Responsible usage and safeguards are essential to prevent potential misuse, such as financial misinformation or unethical market influence[7].

Ethical Statement: The authors take full responsibility for any potential legal issues arising from FinBen's development and dissemination. All data used are publicly available, non-personal, and shared under the MIT license, adhering to privacy and ethical guidelines. This manuscript and associated materials are for academic and educational use only and do not provide financial, legal, or investment advice. The authors disclaim any liability for losses or damages from using the material, and users agree to seek professional consultation and indemnify the authors against any claims arising from its use[8].

---

[7]For a detailed limitation concerning this work, please see Appendix.

[8]For a detailed ethical and legal statement concerning this work, please see Appendix.

References

Yaser S Abu-Mostafa and Amir F Atiya. 1996. Introduction to financial forecasting. *Applied intelligence* 6 (1996), 205–213.

Ebtesam Almazrouei, Hamza Alobeidli, Abdulaziz Alshamsi, Alessandro Cappelli, Ruxandra Cojocaru, Mérouane Debbah, Étienne Goffinet, Daniel Hesslow, Julien Launay, Quentin Malartic, et al. 2023. The falcon series of open language models. *arXiv preprint arXiv:2311.16867* (2023).

Julio Cesar Salinas Alvarado, Karin Verspoor, and Timothy Baldwin. 2015. Domain adaption of named entity recognition to support credit risk assessment. In *Proceedings of the Australasian Language Technology Association Workshop 2015*. 84–90.

Dogu Araci. 2019. FinBERT: Financial Sentiment Analysis with Pre-trained Language Models. arXiv:1908.10063 [cs.CL]

Robert A Ariel. 1987. A monthly effect in stock returns. *Journal of financial economics* 18, 1 (1987), 161–174.

Saqib Aziz and Michael Dowling. 2019. *Machine learning and AI for risk management*. Springer International Publishing.

Baichuan. 2023. Baichuan 2: Open Large-scale Language Models. *arXiv preprint arXiv:2309.10305* (2023). https://arxiv.org/abs/2309.10305

Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in neural information processing systems* 33 (2020), 1877–1901.

Sébastien Bubeck, Varun Chandrasekaran, Ronen Eldan, Johannes Gehrke, Eric Horvitz, Ece Kamar, Peter Lee, Yin Tat Lee, Yuanzhi Li, Scott Lundberg, et al. 2023. Sparks of artificial general intelligence: Early experiments with gpt-4. *arXiv preprint arXiv:2303.12712* (2023).

Longbing Cao. 2022. Ai in finance: challenges, techniques, and opportunities. *ACM Computing Surveys (CSUR)* 55, 3 (2022), 1–38.

Chung-Chi Chen, Yu-Min Tseng, Juyeon Kang, Anaïs Lhuissier, Min-Yuh Day, Teng-Tsai Tu, and Hsin-Hsi Chen. 2023a. Multi-Lingual ESG Issue Identification. In *Proceedings of the Fifth Workshop on Financial Technology and Natural Language Processing and the Second Multimodal AI For Financial Forecasting*. 111–115.

Wei Chen, Qiushi Wang, Zefei Long, Xianyin Zhang, Zhongtian Lu, Bingxuan Li, Siyuan Wang, Jiarong Xu, Xiang Bai, Xuanjing Huang, et al. 2023b. Disc-finllm: A chinese financial large language model based on multiple experts fine-tuning. *arXiv preprint arXiv:2310.15205* (2023).

Zhiyu Chen, Wenhu Chen, Charese Smiley, and et al. Sameena Shah. 2022a. FinQA: A Dataset of Numerical Reasoning over Financial Data. arXiv:2109.00122 [cs.CL]

Zhiyu Chen, Wenhu Chen, Charese Smiley, Sameena Shah, Iana Borova, Dylan Langdon, Reema Moussa, Matt Beane, Ting-Hao Huang, Bryan R Routledge, et al. 2021. FinQA: A Dataset of Numerical Reasoning over Financial Data. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*. 3697–3711.

Zhiyu Chen, Shiyang Li, Charese Smiley, Zhiqiang Ma, Sameena Shah, and William Yang Wang. 2022b. ConvFinQA: Exploring the Chain of Numerical Reasoning in Conversational Finance Question Answering. arXiv:2210.03849 [cs.CL]

Davide Chicco and Giuseppe Jurman. 2020. The advantages of the Matthews correlation coefficient (MCC) over F1 score and accuracy in binary classification evaluation. *BMC genomics* 21, 1 (2020), 1–13.

Keith Cortis, André Freitas, Tobias Daudert, Manuela Huerlimann, Manel Zarrouk, Siegfried Handschuh, and Brian Davis. 2017. Semeval-2017 task 5: Fine-grained sentiment analysis on financial microblogs and news. In *Proceedings of the 11th international workshop on semantic evaluation (SemEval-2017)*. 519–535.

Marco Costantino and Paolo Coletti. 2008. *Information extraction in finance*. Vol. 8. Wit Press.

Yongfu Dai, Duanyu Feng, Jimin Huang, Haochen Jia, Qianqian Xie, Yifang Zhang, Weiguang Han, Wei Tian, and Hao Wang. 2024. LAiW: A Chinese Legal Large Language Models Benchmark. arXiv:2310.05620 [cs.CL]

Leon Derczynski. 2016. Complementarity, F-score, and NLP Evaluation. In *Proceedings of the Tenth International Conference on Language Resources and Evaluation (LREC'16)*, Nicoletta Calzolari, Khalid Choukri, Thierry Declerck, Sara Goggi, Marko Grobelnik, Bente Maegaard, Joseph Mariani, Helene Mazo, Asuncion Moreno, Jan Odijk, and Stelios Piperidis (Eds.). European Language Resources Association (ELRA), Portorož, Slovenia, 261–266. `https://aclanthology.org/L16-1040`

Zhengxiao Du, Yujie Qian, Xiao Liu, Ming Ding, Jiezhong Qiu, Zhilin Yang, and Jie Tang. 2022. GLM: General Language Model Pretraining with Autoregressive Blank Infilling. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 320–335.

Duanyu Feng, Yongfu Dai, Jimin Huang, Yifang Zhang, Qianqian Xie, Weiguang Han, Alejandro Lopez-Lira, and Hao Wang. 2023. Empowering many, biasing a few: Generalist credit scoring through large language models. *arXiv preprint arXiv:2310.00566* (2023).

Cyril Goutte and Eric Gaussier. 2005. A probabilistic interpretation of precision, recall and F-score, with implication for evaluation. In *European conference on information retrieval*. Springer, 345–359.

Weiguang Han, Jimin Huang, Qianqian Xie, Boyi Zhang, Yanzhao Lai, and Min Peng. 2023a. Mastering Pair Trading with Risk-Aware Recurrent Reinforcement Learning. arXiv:2304.00364 [q-fin.CP]

Weiguang Han, Boyi Zhang, Qianqian Xie, Min Peng, Yanzhao Lai, and Jimin Huang. 2023b. Select and Trade: Towards Unified Pair Trading with Hierarchical Reinforcement Learning. *arXiv preprint arXiv:2301.10724* (2023).

Hans Hofmann. 1994. Statlog (German Credit Data). UCI Machine Learning Repository. DOI: https://doi.org/10.24432/C5NC77.

Dong Hongyuan, Che Wanxiang, He Xiaoyu, Zheng Guidong, and Wen Junjie. 2023. FinBART: A Pre-trained Seq2seq Language Model for Chinese Financial Tasks. In *Proceedings of the 22nd Chinese National Conference on Computational Linguistics*, Maosong Sun, Bing Qin, Xipeng Qiu, Jing Jiang, and Xianpei Han (Eds.). Chinese Information Processing Society of China, Harbin, China, 906–917. `https://aclanthology.org/2023.ccl-1.77`

Pranab Islam, Anand Kannappan, Douwe Kiela, Rebecca Qian, Nino Scherrer, and Bertie Vidgen. 2023. FinanceBench: A New Benchmark for Financial Question Answering. *arXiv preprint arXiv:2311.11944* (2023).

Albert Q Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, et al. 2024. Mixtral of experts. *arXiv preprint arXiv:2401.04088* (2024).

Rasmus Jørgensen, Oliver Brandt, Mareike Hartmann, Xiang Dai, Christian Igel, and Desmond Elliott. 2023. MultiFin: A Dataset for Multilingual Financial NLP. In *Findings of the Association for Computational Linguistics: EACL 2023*. 864–879.

Kisub Kim, Xin Zhou, Dongsun Kim, Julia Lawall, Kui Liu, Tegawendé F Bissyandé, Jacques Klein, Jaekwon Lee, and David Lo. 2023. How are We Detecting Inconsistent Method Names? An Empirical Study from Code Review Perspective. *arXiv preprint arXiv:2308.12701* (2023).

Rik Koncel-Kedziorski, Michael Krumdick, Viet Lai, Varshini Reddy, Charles Lovering, and Chris Tanner. 2023. Bizbench: A quantitative reasoning benchmark for business and finance. *arXiv preprint arXiv:2311.06602* (2023).

Moreno La Quatra and Luca Cagliero. 2020. End-to-end training for financial report summarization. In *Proceedings of the 1st Joint Workshop on Financial Narrative Processing and MultiLing Financial Summarisation*. 118–123.

Matthew Lamm, Arun Tejasvi Chaganty, Christopher D Manning, Dan Jurafsky, and Percy Liang. 2018. Textual analogy parsing: What's shared and what's compared among analogous facts. *arXiv preprint arXiv:1809.02700* (2018).

Jean Lee, Nicholas Stevens, Soyeon Caren Han, and Minseok Song. 2024. A Survey of Large Language Models in Finance (FinLLMs). arXiv:2402.02315 [cs.CL]

Yang Lei, Jiangtong Li, Ming Jiang, Junjie Hu, Dawei Cheng, Zhijun Ding, and Changjun Jiang. 2023. CFBenchmark: Chinese Financial Assistant Benchmark for Large Language Model. arXiv:2311.05812 [cs.CL]

Jiangtong Li, Yuxuan Bian, Guoxuan Wang, Yang Lei, Dawei Cheng, Zhijun Ding, and Changjun Jiang. 2023a. CFGPT: Chinese Financial Assistant with Large Language Model. arXiv:2309.10654 [cs.CL]

Xianzhi Li, Xiaodan Zhu, Zhiqiang Ma, Xiaomo Liu, and Sameena Shah. 2023c. Are ChatGPT and GPT-4 General-Purpose Solvers for Financial Text Analytics? An Examination on Several Typical Tasks. *arXiv preprint arXiv:2305.05862* (2023).

Yinheng Li, Shaofei Wang, Han Ding, and Hang Chen. 2023b. Large Language Models in Finance: A Survey. arXiv:2311.10723 [q-fin.GN]

Chin-Yew Lin. 2004. Rouge: A package for automatic evaluation of summaries. In *Text summarization branches out*. 74–81.

Xiao-Yang Liu, Guoxuan Wang, and Daochen Zha. 2023a. Fingpt: Democratizing internet-scale data for financial large language models. *arXiv preprint arXiv:2307.10485* (2023).

Xiao-Yang Liu, Ziyi Xia, Jingyang Rui, Jiechao Gao, Hongyang Yang, Ming Zhu, Christina Dan Wang, Zhaoran Wang, and Jian Guo. 2022. FinRL-Meta: Market Environments and Benchmarks for Data-Driven Financial Reinforcement Learning. arXiv:2211.03107 [q-fin.TR]

Xiao-Yang Liu, Ziyi Xia, Hongyang Yang, Jiechao Gao, Daochen Zha, Ming Zhu, Christina Dan Wang, Zhaoran Wang, and Jian Guo. 2023b. Dynamic Datasets and Market Environments for Financial Reinforcement Learning. arXiv:2304.13174 [cs.LG]

Zhuang Liu, Degen Huang, Kaiyu Huang, Zhuang Li, and Jun Zhao. 2020. FinBERT: A Pre-trained Financial Language Representation Model for Financial Text Mining. In *Proceedings of the Twenty-Ninth International Joint Conference on Artificial Intelligence, IJCAI-20*, Christian Bessiere (Ed.). International Joint Conferences on Artificial Intelligence Organization, 4513–4519. https://doi.org/10.24963/ijcai.2020/622 Special Track on AI in FinTech.

Alejandro Lopez-Lira and Yuehua Tang. 2023. Can chatgpt forecast stock price movements? return predictability and large language models. *arXiv preprint arXiv:2304.07619* (2023).

Tim Loughran and Bill McDonald. 2020. Textual analysis in finance. *Annual Review of Financial Economics* 12 (2020), 357–375.

Malik Magdon-Ismail and Amir F Atiya. 2004. Maximum drawdown. *Risk Magazine* 17, 10 (2004), 99–102.

Macedo Maia, Siegfried Handschuh, Andre Freitas, Brian Davis, Ross McDermott, Manel Zarrouk, and Alexandra Balahur. 2018. WWW'18 Open Challenge: Financial Opinion Mining and Question Answering. *WWW '18: Companion Proceedings of the The Web Conference 2018*, 1941–1942. https://doi.org/10.1145/3184558.3192301

Pekka Malo, Ankur Sinha, Pekka Korhonen, Jyrki Wallenius, and Pyry Takala. 2014. Good debt or bad debt: Detecting semantic orientations in economic texts. *Journal of the Association for Information Science and Technology* 65, 4 (2014), 782–796.

Dominique Mariko, Hanna Abi Akl, Estelle Labidurie, Stephane Durfort, Hugues De Mazancourt, and Mahmoud El-Haj. 2020. Financial document causality detection shared task (fincausal 2020). *arXiv preprint arXiv:2012.02505* (2020).

Rajdeep Mukherjee, Abhinav Bohra, Akash Banerjee, Soumya Sharma, Manjunath Hegde, Afreen Shaikh, Shivani Shrivastava, Koustuv Dasgupta, Niloy Ganguly, Saptarshi Ghosh, et al. 2022. Ectsum: A new benchmark dataset for bullet point summarization of long earnings call transcripts. *arXiv preprint arXiv:2210.12467* (2022).

OpenAI. 2023a. GPT-4 Technical Report. arXiv:2303.08774 [cs.CL]

R OpenAI. 2023b. Gpt-4 technical report. arxiv 2303.08774. *View in Article* 2 (2023), 13.

Felipe Dias Paiva, Rodrigo Tomás Nogueira Cardoso, Gustavo Peixoto Hanaoka, and Wendel Moreira Duarte. 2019. Decision-making for financial trading: A fusion approach of machine learning and portfolio selection. *Expert Systems with Applications* 115 (2019), 635–655.

André E Punt. 2017. Strategic management decision-making in a complex world: quantifying, understanding, and using trade-offs. *ICES Journal of Marine Science* 74, 2 (2017), 499–510.

Ross Quinlan. [n. d.]. Statlog (Australian Credit Approval). UCI Machine Learning Repository. DOI: https://doi.org/10.24432/C59012.

Baptiste Roziere, Jonas Gehring, Fabian Gloeckle, Sten Sootla, Itai Gat, Xiaoqing Ellen Tan, Yossi Adi, Jingyu Liu, Tal Remez, Jérémy Rapin, et al. 2023. Code llama: Open foundation models for code. *arXiv preprint arXiv:2308.12950* (2023).

Julio Cesar Salinas Alvarado, Karin Verspoor, and Timothy Baldwin. 2015. Domain Adaption of Named Entity Recognition to Support Credit Risk Assessment. In *Proceedings of the Australasian Language Technology Association Workshop 2015*, Ben Hachey and Kellie Webster (Eds.). Parramatta, Australia, 84–90. https://aclanthology.org/U15-1010

Agam Shah, Suvan Paturi, and Sudheer Chava. 2023a. Trillion Dollar Words: A New Financial Dataset, Task & Market Analysis. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, Anna Rogers, Jordan Boyd-Graber, and Naoaki Okazaki (Eds.). Association for Computational Linguistics, Toronto, Canada, 6664–6679. https://doi.org/10.18653/v1/2023.acl-long.368

Agam Shah, Ruchit Vithani, Abhinav Gullapalli, and Sudheer Chava. 2023b. Finer: Financial named entity recognition dataset and weak-supervision model. *arXiv preprint arXiv:2302.11157* (2023).

Raj Shah, Kunal Chawla, Dheeraj Eidnani, Agam Shah, Wendi Du, Sudheer Chava, Natraj Raman, Charese Smiley, Jiaao Chen, and Diyi Yang. 2022. When FLUE Meets FLANG: Benchmarks and Large Pretrained Language Model for Financial Domain. In *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing*. 2322–2335.

Soumya Sharma, Subhendu Khatuya, Manjunath Hegde, Afreen Shaikh, Koustuv Dasgupta, Pawan Goyal, and Niloy Ganguly. 2023. Financial Numeric Extreme Labelling: A dataset and benchmarking. In *Findings of the Association for Computational Linguistics: ACL 2023*. 3550–3561.

Soumya Sharma, Tapas Nayak, Arusarka Bose, Ajay Kumar Meena, Koustuv Dasgupta, Niloy Ganguly, and Pawan Goyal. 2022. FinRED: A dataset for relation extraction in financial domain. In *Companion Proceedings of the Web Conference 2022*. 595–597.

William F Sharpe. 1998. The sharpe ratio. *Streetwise–the Best of the Journal of Portfolio Management* 3 (1998), 169–85.

Ankur Sinha and Tanmay Khandait. 2020. Impact of News on the Commodity Market: Dataset and Results. arXiv:2009.04202 [cs.CL]

Ankur Sinha and Tanmay Khandait. 2021. Impact of news on the commodity market: Dataset and results. In *Advances in Information and Communication: Proceedings of the 2021 Future of Information and Communication Conference (FICC), Volume 2*. Springer, 589–601.

Yejun Soun, Jaemin Yoo, Minyong Cho, Jihyeong Jeon, and U Kang. 2022. Accurate Stock Movement Prediction with Self-supervised Learning from Sparse Noisy Tweets. In *2022 IEEE International Conference on Big Data (Big Data)*. IEEE, 1691–1700.

Eugene Sy, Tzu-Cheng Peng, Shih-Hsuan Huang, Heng-Yu Lin, and Yung-Chun Chang. 2023. Fine-Grained Argument Understanding with BERT Ensemble Techniques: A Deep Dive into Financial Sentiment Analysis. In *Proceedings of the 35th Conference on Computational Linguistics and Speech Processing (ROCLING 2023)*. 242–249.

Gemini Team, Rohan Anil, Sebastian Borgeaud, Yonghui Wu, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, et al. 2023. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805* (2023).

InternLM Team. 2023. Internlm: A multilingual language model with progressively enhanced capabilities.

Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, et al. 2023a. Llama: Open and efficient foundation language models. *arXiv preprint arXiv:2302.13971* (2023).

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. 2023b. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288* (2023).

Neng Wang, Hongyang Yang, and Christina Dan Wang. 2023. FinGPT: Instruction Tuning Benchmark for Open-Source Large Language Models in Financial Datasets. arXiv:2310.04793 [cs.CL]

Huizhe Wu, Wei Zhang, Weiwei Shen, and Jun Wang. 2018. Hybrid deep sequential modeling for social text-driven stock prediction. In *Proceedings of the 27th ACM international conference on information and knowledge management*. 1627–1630.

Shijie Wu, Ozan Irsoy, Steven Lu, Vadim Dabravolski, Mark Dredze, Sebastian Gehrmann, Prabhanjan Kambadur, David Rosenberg, and Gideon Mann. 2023. BloombergGPT: A Large Language Model for Finance. arXiv:2303.17564 [cs.LG]

Qianqian Xie, Weiguang Han, Yanzhao Lai, Min Peng, and Jimin Huang. 2023a. The Wall Street Neophyte: A Zero-Shot Analysis of ChatGPT Over MultiModal Stock Movement Prediction Challenges. *arXiv preprint arXiv:2304.05351* (2023).

Qianqian Xie, Weiguang Han, Xiao Zhang, Yanzhao Lai, Min Peng, Alejandro Lopez-Lira, and Jimin Huang. 2023b. PIXIU: A Large Language Model, Instruction Data and Evaluation Benchmark for Finance. *arXiv preprint arXiv:2306.05443* (2023).

Yumo Xu and Shay B Cohen. 2018. Stock movement prediction from tweets and historical prices. In *Proceedings of the 56th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*. 1970–1979.

Hongyang Yang, Xiao-Yang Liu, and Christina Dan Wang. 2023a. FinGPT: Open-Source Financial Large Language Models. arXiv:2306.06031 [q-fin.ST]

Linyi Yang, Eoin M Kenny, Tin Lok James Ng, Yi Yang, Barry Smyth, and Ruihai Dong. 2020a. Generating plausible counterfactual explanations for deep transformers in financial text classification. *arXiv preprint arXiv:2010.12512* (2020).

Yi Yang, Yixuan Tang, and Kar Yan Tam. 2023b. InvestLM: A Large Language Model for Investment using Financial Domain Instruction Tuning. arXiv:2309.13064 [q-fin.GN]

Yi Yang, Mark Christopher Siy UY, and Allen Huang. 2020b. FinBERT: A Pretrained Language Model for Financial Communications. arXiv:2006.08097 [cs.CL]

Yangyang Yu, Haohang Li, Zhi Chen, Yuechen Jiang, Yang Li, Denghui Zhang, Rong Liu, Jordan W. Suchow, and Khaldoun Khashanah. 2023. FinMem: A Performance-Enhanced LLM Trading Agent with Layered Memory and Character Design. arXiv:2311.13743 [q-fin.CP]

Weizhe Yuan, Graham Neubig, and Pengfei Liu. 2021. Bartscore: Evaluating generated text as text generation. *Advances in Neural Information Processing Systems* 34 (2021), 27263–27277.

Shengyu Zhang, Linfeng Dong, Xiaoya Li, Sen Zhang, Xiaofei Sun, Shuhe Wang, Jiwei Li, Runyi Hu, Tianwei Zhang, Fei Wu, and Guoyin Wang. 2023a. Instruction Tuning for Large Language Models: A Survey. arXiv:2308.10792 [cs.CL]

Tianyi Zhang, Varsha Kishore, Felix Wu, Kilian Q Weinberger, and Yoav Artzi. 2019. Bertscore: Evaluating text generation with bert. *arXiv preprint arXiv:1904.09675* (2019).

Xuanyu Zhang, Bingbing Li, and Qing Yang. 2023b. CGCE: A Chinese Generative Chat Evaluation Benchmark for General and Financial Domains. arXiv:2305.14471 [cs.CL]

Xiao Zhang, Ruoyu Xiang, Chenhan Yuan, Duanyu Feng, Weiguang Han, Alejandro Lopez-Lira, Xiao-Yang Liu, Sophia Ananiadou, Min Peng, Jimin Huang, and Qianqian Xie. 2024. Dólares or Dollars? Unraveling the Bilingual Prowess of Financial LLMs Between Spanish and English. arXiv:2402.07405 [cs.CL]

Xuanyu Zhang, Qing Yang, and Dongliang Xu. 2023c. XuanYuan 2.0: A Large Chinese Financial Chat Model with Hundreds of Billions Parameters. arXiv:2305.12002 [cs.CL]

Huaqin Zhao, Zhengliang Liu, Zihao Wu, Yiwei Li, Tianze Yang, Peng Shu, Shaochen Xu, Haixing Dai, Lin Zhao, Gengchen Mai, et al. 2024. Revolutionizing finance with llms: An overview of applications and insights. *arXiv preprint arXiv:2401.11641* (2024).

Xianzheng Zhou, Hui Zhou, and Huaigang Long. 2023. Forecasting the equity premium: Do deep neural network models work? *Modern Finance* 1, 1 (2023), 1–11.

Zhihan Zhou, Liqian Ma, and Han Liu. 2021. Trade the Event: Corporate Events Detection for News-Based Event-Driven Trading. arXiv:2105.12825 [cs.CL]

Fengbin Zhu, Wenqiang Lei, Youcheng Huang, Chao Wang, Shuo Zhang, Jiancheng Lv, Fuli Feng, and Tat-Seng Chua. 2021. TAT-QA: A question answering benchmark on a hybrid of tabular and textual content in finance. *arXiv preprint arXiv:2105.07624* (2021).

Checklist

1. For all authors...
   (a) Do the main claims made in the abstract and introduction accurately reflect the paper's contributions and scope? [Yes]
   (b) Did you describe the limitations of your work? [Yes] See Limitation (Section G.2).
   (c) Did you discuss any potential negative societal impacts of your work? [Yes] See Ethical Statement (Section G.2).
   (d) Have you read the ethics review guidelines and ensured that your paper conforms to them? [Yes] See Ethical Statement (Section G.2).

2. If you are including theoretical results...
   (a) Did you state the full set of assumptions of all theoretical results? [N/A]
   (b) Did you include complete proofs of all theoretical results? [N/A]

3. If you ran experiments (e.g. for benchmarks)...
   (a) Did you include the code, data, and instructions needed to reproduce the main experimental results (either in the supplemental material or as a URL)? [Yes] See Introduction (Section 1).
   (b) Did you specify all the training details (e.g., data splits, hyperparameters, how they were chosen)? [N/A] Our benchmark only includes the evaluation process.
   (c) Did you report error bars (e.g., with respect to the random seed after running experiments multiple times)? [Yes] See Table 4.
   (d) Did you include the total amount of compute and the type of resources used (e.g., type of GPUs, internal cluster, or cloud provider)? [Yes] See Experimental Settings (Section 3).

4. If you are using existing assets (e.g., code, data, models) or curating/releasing new assets...

   (a) If your work uses existing assets, did you cite the creators? [Yes]
   (b) Did you mention the license of the assets? [Yes]
   (c) Did you include any new assets either in the supplemental material or as a URL? [Yes] Our Introduction (Section 1) contains a link for all data used in FinBen.
   (d) Did you discuss whether and how consent was obtained from people whose data you're using/curating? [Yes] Our dataset statistics (Table 2) contains licenses for all used datasets.
   (e) Did you discuss whether the data you are using/curating contains personally identifiable information or offensive content? [Yes] See Ethical Statement (Section G.2).

5. If you used crowdsourcing or conducted research with human subjects...

   (a) Did you include the full text of instructions given to participants and screenshots, if applicable? [N/A]
   (b) Did you describe any potential participant risks, with links to Institutional Review Board (IRB) approvals, if applicable? [N/A]
   (c) Did you include the estimated hourly wage paid to participants and the total amount spent on participant compensation? [N/A]

A Contributions

Science Leadership: Qianqian Xie, Min Peng, Sophia Ananiadou, Alejandro Lopez-Lira, Hao Wang, Yanzhao Lai, Benyou Wang, Xiao-yang Liu, Gang Hu, Jiajia Huang, Jimin Huang.

Contributors: Mengxi Xiao, Dong Li, Weiguang Han, Zhengyu Chen, Ruoyu Xiang, Xiao Zhang, Yueru He, Yongfu Dai, Duanyu Feng, Yijing Xu, Haoqiang Kang, Ziyan Kuang, Chenhan Yuan, Kailai Yang, Zheheng Luo, Tianlin Zhang, Zhiwei Liu, Guojun Xiong, Zhiyang Deng, Yuechen Jiang, Zhiyuan Yao, Haohang Li, Yangyang Yu

B Other LLMs Performance

Table 6 presents other LLMs' performance in the FinBen.

C Instructions

For detail instruction of each dataset, please see Table 7 and Table 8.

D Related Work

D.1 Financial Large Language Models

Recent years have seen a significant surge in research on finance-specific LLMs, expanding on the groundwork laid by general-purpose language models (Lee et al., 2024; Liu et al., 2023b; Xie et al., 2023a; Zhang et al., 2024; Dai et al., 2024). Financial pre-trained language models (FinPLMs) like FinBERT (Araci, 2019; Yang et al., 2020b; Liu et al., 2020), derived from BERT, and FLANG (Shah et al., 2022), based on ELECTRA, have been developed using domain-specific data for enhanced performance in tasks like sentiment analysis and stock prediction. The open-source release of Meta AI's LLaMA (Touvron et al., 2023a,b) has fueled further innovation in Financial LLMs (FinLLMs), with models like FinMA (Xie et al., 2023b), InvestLM (Yang et al., 2023b), and FinGPT (Wang et al., 2023; Liu et al., 2023a) leveraging advanced tuning strategies (Zhang et al., 2023a) for financial applications. BloombergGPT (Wu et al., 2023) stands out as a BLOOM-based, closed-source model tailored for the financial industry. Additionally, the Chinese financial sector has seen the emergence of models like XuanYuan 2.0 (Zhang et al., 2023c), integrating broad and specialized knowledge, FinBART (Hongyuan et al., 2023) for financial communication, and CFGPT (Li et al., 2023a), which includes a comprehensive dataset for targeted pre-training and fine-tuning.

Table 6: The zero-shot and few-shots performance of other LLMs on the FinBen.

Dataset	Metrics	Baichuan 7B	CodeLLaMA 7B	DISC-FinLLM	ChatGLM3 6B
NER	EntityF1	0.00	0.07	0.12	**0.25**
FINER-ORD	EntityF1	0.00	0.00	0.00	**0.02**
FinRED	F1	**0.00**	0.00	0.00	0.00
SC	F1	0.74	**0.85**	0.00	0.81
CD	F1	**0.00**	0.00	0.00	0.00
FNXL	EntityF1	**0.00**	0.00	0.00	0.00
FSRL	EntityF1	**0.00**	0.00	0.00	0.00
FPB	F1	0.17	0.34	0.29	**0.74**
	Acc	0.23	0.39	0.26	**0.74**
FiQA-SA	F1	0.32	**0.66**	0.32	0.56
TSA	RMSE↓	0.44	0.43	**0.32**	0.35
Headlines	AvgF1	0.60	0.60	0.60	**0.66**
FOMC	F1	0.17	0.14	0.19	**0.47**
	Acc	0.25	0.27	0.28	**0.46**
FinArg-ACC	MicroF1	**0.36**	0.28	0.29	0.25
FinArg-ARC	MicroF1	0.27	0.25	0.29	**0.50**
MultiFin	MicroF1	0.12	0.21	0.29	**0.47**
M&A	MicroF1	0.33	0.54	0.29	**0.79**
MLESG	MicroF1	0.04	0.10	**0.29**	0.16
FinQA	EmAcc	**0.00**	0.00	0.00	0.00
TATQA	EmAcc	0.00	0.00	0.00	**0.07**
Regulations	Rouge-1	0.13	-	-	**0.26**
	BertScore	0.60	-	-	**0.65**
ConvFinQA	EmAcc	**0.00**	0.00	0.00	0.00
EDTSUM	Rouge-1	0.02	0.10	**0.22**	0.13
	BertScore	0.47	**0.67**	0.61	0.47
	BartScore	-6.18	**-3.62**	-4.13	-4.78
ECTSUM	Rouge-1	**0.00**	0.00	0.00	0.00
	BertScore	**0.00**	0.00	0.00	0.00
	BartScore	**-5.18**	-5.18	-5.18	-5.18
BigData22	Acc	**0.53**	0.52	0.44	0.47
	MCC	-0.01	-0.01	-0.05	**0.00**
ACL18	Acc	0.50	**0.51**	0.50	0.50
	MCC	-0.01	0.00	**0.02**	0.02
CIKM18	Acc	0.48	**0.51**	0.44	0.42
	MCC	**0.02**	0.02	-0.03	0.02
German	F1	0.52	**0.66**	0.52	0.41
	MCC	**0.00**	0.00	0.00	-0.30
Australian	F1	0.26	**0.43**	0.26	0.27
	MCC	**0.00**	0.00	0.00	-0.02
LendingClub	F1	0.59	**0.81**	0.72	0.72
	MCC	**0.02**	0.00	0.00	-0.03
ccf	F1	0.01	0.00	0.66	**1.00**
	MCC	**0.00**	0.00	-0.04	0.00
ccfraud	F1	0.05	0.06	0.46	**0.67**
	MCC	-0.04	0.00	**0.02**	-0.15
polish	F1	**0.92**	0.47	0.92	0.05
	MCC	-0.01	**0.04**	0.00	0.00
taiwan	F1	**0.95**	0.36	0.95	0.05
	MCC	0.00	-0.03	0.00	**0.01**
portoseguro	F1	0.00	0.88	0.63	**0.95**
	MCC	0.00	-0.01	-0.02	**0.06**
travelinsurance	F1	0.00	0.02	0.00	**0.97**
	MCC	0.00	0.00	0.00	**0.03**

D.2 Financial Evaluation Benchmarks

Financial evaluation benchmarks, such as the pioneering FLUE (Shah et al., 2022), have been introduced to measure model performance in the financial sector, covering five key NLP tasks: financial sentiment analysis (Shah et al., 2022), news headline classification (Sinha and Khandait, 2020), named entity recognition (NER) (Salinas Alvarado et al., 2015), structure boundary detection and question answering (QA) (Chen et al., 2022a). Building upon FLUE, FLARE (Xie et al., 2023b) added the evaluation of time-series processing capabilities, i.e., forecasting stock price movements. In addition, in Chinese financial benchmarks, there are more recently released Chinese datasets like CFBenchmark (Lei et al., 2023), DISC-FINSFT (Chen et al., 2023b), and CGCE (Zhang et al., 2023b). However, these benchmarks have a limited scope and have not yet addressed more complex

Table 7: Quantification task datasets prompt overview.

Data	Prompt
FPB	"Analyze the sentiment of this statement extracted from a financial news article. Provide your answer as either negative, positive or neutral. For instance, 'The company's stocks plummeted following the scandal.' would be classified as negative."
FiQA-SA	"What is the sentiment of the following financial {category}: Positive, Negative, or Neutral?"
Headlines	"Consider whether the headline mentions the price of gold. Is there a Price or Not in the gold commodity market indicated in the news headline? Please answer Yes or No."
NER	"In the sentences extracted from financial agreements in U.S. SEC filings, identify the named entities that represent a person ('PER'), an organization ('ORG'), or a location ('LOC'). The required answer format is: 'entity name, entity type'. For instance, in 'Elon Musk, CEO of SpaceX, announced the launch from Cape Canaveral.', the entities would be: 'Elon Musk, PER; SpaceX, ORG; Cape Canaveral, LOC'"
FiNER-ORD	"In the list of tokens, identify {tid}each accordingly. If the entity spans multiple tokens, use the prefix B-PER, B-LOC, or B-ORG for the first token, and I-PER, I-LOC, or I-ORG for the subsequent tokens of that entity. The beginning of each separate entity should always be labeled with a B-PER, B-LOC, or B-ORG prefix. If the token does not fit into any of the three named categories, or is not a named entity, label it as 'O'."
FinQA	"Given the financial data and expert analysis, please answer this question:"
Regulations	"Please answer following questions."
ConvFinQA	"In the context of this series of interconnected finance-related queries and the additional information provided by the pretext, table data, and post text from a company's financial filings, please provide a response to the final question. This may require extracting information from the context and performing mathematical calculations. Please take into account the information provided in the preceding questions and their answers when formulating your response:"
BigData22	" Contemplate the data and tweets to guess whether the closing price of {tid} will surge or decline at {point}. Please declare with either Rise or Fall."
ACL18	"Scrutinize the data and tweets to envisage if the closing price of {tid}will swell or contract at {point}. Respond with either Rise or Fall."
CIKM18	"Reflect on the provided data and tweets to anticipate if the closing price of {tid}is going to increase or decrease at {point}. Respond with either Rise or Fall."
ECTSum	"Given the following article, please produce a list of 0 and 1, each separated by ' ' to indicate which sentences should be included in the final summary. The article's sentences have been split by ' '. Please mark each sentence with 1 if it should be included in the summary and 0 if it should not."
EDTSum	"You are given a text that consists of multiple sentences. Your task is to perform abstractive summarization on this text. Use your understanding of the content to express the main ideas and crucial details in a shorter, coherent, and natural sounding text."
German	"Assess the creditworthiness of a customer using the following table attributes for financial status. Respond with either 'good' or 'bad'. And the table attributes including 13 categorical attributes and 7 numerical attributes are as follows:"
Australian	"Assess the creditworthiness of a customer using the following table attributes for financial status. Respond with either 'good' or 'bad'. And the table attributes including 13 categorical attributes and 7 numerical attributes and values have been changed to meaningless symbols to protect confidentiality of the data. :"
FOMC	"Examine the excerpt from a central bank's release below. Classify it as HAWKISH if it advocates for a tightening of monetary policy, DOVISH if it suggests an easing of monetary policy, or NEUTRAL if the stance is unbiased. Your response should return only HAWKISH, DOVISH, or NEUTRAL."
TSA	"Given the following financial text, return a sentiment score for Ashtead as a floating-point number ranging from -1 (indicating a very negative or bearish sentiment) to 1 (indicating a very positive or bullish sentiment), with 0 designating neutral sentiment. Return only the numerical score first, follow it with a brief reasoning behind your score."
FinArg - ACC	"Analyze sentences from earnings conference calls and identify their argumentative function. Each sentence is either a premise, offering evidence or reasoning, or a claim, asserting a conclusion or viewpoint. Return only premise or claim."
FinArg - ARC	"In this task, you are given a pair of sentences. Your objective is to ascertain the type of argumentative relation between these two sentences. The relation could either be 'NoRelation', indicating no discernible relation between the sentences, 'Support', indicating that the first sentence supports the second, or 'Attack', indicating that the first sentence disputes or contradicts the second. Return only one of the three classifications: 'norelation', 'support', or 'attack'."
MultiFin	"In this task, you're working with English headlines from the MULTIFIN dataset. This dataset is made up of real-world article headlines from a large accounting firm's websites. Your objective is to categorize each headline according to its primary topic. The potential categories are {category}. Your response should only include the category that best fits the headline."
MA	"In this task, you will be given Mergers and Acquisitions news articles or tweets. Your task is to classify each article or tweet based on whether the mentioned deal was completed or remained a rumour. Your response should be a single word - either 'complete' or 'rumour' - representing the outcome of the deal mentioned in the provided text."
MLESG	"You're given English news articles related to Environmental, Social, and Corporate Governance (ESG) issues. Your task is to classify each article based on the ESG issue it pertains to, according to the MSCI ESG rating guidelines. The ESG issues include {category}. Your output should be the most relevant ESG issue label, followed by a brief rationale based on the article content."

financial NLP tasks such as event detection (Zhou et al., 2021), and realistic financial tasks, despite the fact that there were previous efforts on stock trading (Liu et al., 2022; Han et al., 2023a,b).

Table 8: The example prompts of remaining tasks. FiQA-SA has two types of text, including news headlines and tweets. We will fill the detailed text type into {category} for each data sample. For stock movement prediction data such as BigData22, we will fill {tid} and {point} with the detailed stock name and time from each data sample.

Data	Prompt
FinRED	"Given the following sentence, identify the head, tail, and relation of each triplet present in the sentence. The relations you should be looking for are {category}. If a relation exists between two entities, provide your answer in the format {category}. If there are multiple triplets in a sentence, provide each one on a new line."
SC	"In this task, you are provided with sentences extracted from financial news and SEC data. Your goal is to classify each sentence into either 'causal' or 'noise' based on whether or not it indicates a causal relationship between financial events. Please return only the category 'causal' or 'noise'."
CD	"Your job in this task is to perform sequence labeling on a provided text section, marking the chunks that represent the cause of an event and the effects that result from it. For each token in the text, assign a label to indicate its role in representing cause or effect. The labels you should use are 'B-CAUSE', 'I-CAUSE', 'B-EFFECT', 'I-EFFECT', and 'O'. A 'B-' prefix is used to denote the beginning of a cause or effect sequence, while an 'I-' prefix is used for continuation of a cause or effect sequence. If a token is not part of either a cause or effect sequence, label it as 'O'. Provide your answer as a sequence of 'token:label' pairs, with each pair on a new line."
TATQA	"Please answer the given financial question based on the context. Context: {context}Question: What is the amount of total sales in 2019?"
FNXL	"In the task of Financial Numeric Extreme Labelling (FNXL), your job is to identify and label the semantic role of each token in a sentence. The labels can include {category}"
FSRL	"In the task of Textual Analogy Parsing (TAP), your job is to identify and label the semantic role of each token in a sentence. The labels can include {category}."
LendingClub	"Assess the client's loan status based on the following loan records from Lending Club. Respond with only 'good' or 'bad', and do not provide any additional information. For instance, 'The client has a stable income, no previous debts, and owns a property.' should be classified as 'good'."
ccf	"Detect the credit card fraud using the following financial table attributes. Respond with only 'yes' or 'no', and do not provide any additional information. Therein, the data contains 28 numerical input variables V1, V2, ..., and V28 which are the result of a PCA transformation and 1 input variable Amount which has not been transformed with PCA. The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning. For instance, 'The client has attributes:{category}'"
ccfraud	"Detect the credit card fraud with the following financial profile. Respond with only 'good' or 'bad', and do not provide any additional information. For instance, 'The client is a female, the state number is 25, the number of cards is 1, the credit balance is 7000, the number of transactions is 16, the number of international transactions is 0, the credit limit is 6.' should be classified as 'good'."
polish	"Predict whether the company will face bankruptcy based on the financial profile attributes provided in the following text. Respond with only 'no' or 'yes', and do not provide any additional information."
taiwan	"Predict whether the company will face bankruptcy based on the financial profile attributes provided in the following text. Respond with only 'no' or 'yes', and do not provide any additional information."
Porto-Seguro	"Identify whether or not to files a claim for the auto insurance policy holder using the following table attributes about individual financial profile. Respond with only 'yes' or 'no', and do not provide any additional information. And the table attributes that belong to similar groupings are tagged as such in the feature names (e.g., ind, reg, car, calc). In addition, feature names include the postfix bin to indicate binary features and cat to indicate categorical features. Features without these designations are either continuous or ordinal. Values of -1 indicate that the feature was missing from the observation."
travelinsurace	"Identify the claim status of insurance companies using the following table attributes for travel insurance status. Respond with only 'yes' or 'no', and do not provide any additional information. And the table attributes including 5 categorical attributes and 4 numerical attributes are as follows:{category}"
FinTrade	"Given the information, can you make an investment decision? Just summarize the reason of the decision. please consider only the available short-term information, the mid-term information, the long-term information, the reflection-term information. please consider the momentum of the historical stock price. When cumulative return is positive or zero, you are a risk-seeking investor. But when cumulative return is negative, you are a risk-averse investor. please consider how much share of the stock the investor holds now. You should provide exactly one of the following investment decisions: buy or sell. When it is really hard to make a 'buy'-or-'sell' decision, you could go with 'hold' option. You also need to provide the id of the information to support your decision. {investment_info} {grcomplete_json_suffix_v2} Your output should strictly conforms the following json format without any additional contents: {"investment_decision" : string, "summary_reason": string, "short_memory_index": number, "middle_memory_index": number, "long_memory_index": number, "reflection_memory_index": number}"

E Trading Accumulative Returns

Table 9 and the below Figures show detailed trading performance,

Table 9: The overall trading performance comparison for different LLMs across various stocks. The results include large LLMs only ($\geq 70B$), as models with smaller contexts have difficulty understanding the instructions and producing a static strategy of holding.

Ticker	Model	CR (%)	SR	DV (%)	AV (%)	MD (%)
TSLA	Buy & Hold	-25.2137	-0.7203	4.4099	70.0043	57.6765
	GPT-4	**68.3089**	**2.8899**	2.9780	**47.2739**	**10.7996**
	GPT-4o	-0.8789	-0.0321	3.4531	54.8156	44.6842
	GPT3.5-Turbo	25.2137	0.7203	4.4099	70.0043	51.3186
	llama2-70B	-31.4144	-1.0412	3.8014	60.3450	48.6173
	llama3-70B	-16.4424	-0.4847	4.2743	67.8519	55.5486
	gemini	-0.3790	-0.0148	3.2271	51.2280	35.6707
NFLX	Buy & Hold	34.6251	1.3696	3.1852	50.5634	20.9263
	GPT-4	**36.4485**	**2.0088**	2.2860	36.2894	**15.8495**
	GPT-4o	5.5829	0.2592	2.7132	43.0702	17.4715
	GPT3.5-Turbo	7.9337	0.4610	2.1680	34.4160	17.9578
	llama2-70B	33.8460	1.4741	2.8928	45.9216	20.3910
	llama3-70B	21.7374	0.9513	2.8788	45.6989	21.3478
	gemini	11.6298	1.0073	**1.4546**	**23.0906**	16.5106
AMZN	Buy & Hold	-16.4428	-0.7448	2.7812	44.1508	33.8847
	GPT-4	10.5539	0.4923	2.7012	42.8802	22.9294
	GPT-4o	11.3626	0.7334	1.9520	30.9864	19.5964
	GPT3.5-Turbo	**19.9636**	0.9611	2.6171	41.5454	19.2191
	llama2-70B	8.3595	**1.9715**	**0.5342**	**8.4804**	**0.0000**
	llama3-70B	11.1479	0.5405	2.5986	41.2509	28.2174
	gemini	-2.3838	-0.5321	0.5645	8.9605	6.4291
MSFT	Buy & Hold	17.2161	0.9710	2.2339	35.4623	15.0097
	GPT-4	25.7826	**1.5818**	2.0535	32.5989	**14.9889**
	GPT-4o	-5.3731	-0.5209	**1.2997**	**20.6317**	18.8223
	GPT3.5-Turbo	20.4179	1.3600	1.8915	30.0259	20.3212
	llama2-70B	**27.7664**	1.5708	2.2270	35.3524	15.0097
	llama3-70B	21.1983	1.2628	2.1149	33.5724	15.0097
	gemini	21.5081	1.3701	1.9777	31.3957	17.5051
AAPL	Buy & Hold	12.7371	0.7759	2.0682	32.8323	20.6590
	GPT-4	**21.2335**	**1.9274**	1.3879	22.0328	6.4237
	GPT-4o	-6.7540	-0.5693	1.4948	23.7285	20.7600
	GPT3.5-Turbo	0.7110	0.0758	**1.1817**	**18.7581**	**6.0818**
	llama2-70B	11.4856	1.1550	1.2529	19.8885	9.2776
	llama3-70B	-16.0835	-1.1985	1.6907	26.8394	25.9520
	gemini	18.1718	1.7214	1.3300	21.1134	9.6467
GOOG	Buy & Hold	6.3107	0.3081	2.5806	40.9660	21.1907
	GPT-4	13.2811	0.9667	1.7308	27.4762	12.2209
	GPT-4o	16.5072	1.0654	1.9520	30.9872	**11.8863**
	GPT3.5-Turbo	0.9990	0.0614	2.0490	32.5265	20.9316
	llama2-70B	17.0030	**1.1057**	1.9374	30.7546	13.2088
	llama3-70B	17.5630	0.8872	2.4942	39.5934	19.2783
	gemini	**38.7956**	3.0341	**1.6110**	**25.5732**	13.7311
DIS	Buy & Hold	-0.0700	-0.0037	2.3667	37.5695	22.7722
	GPT-4	**31.3383**	**2.3931**	1.6498	26.1904	12.3417
	GPT-4o	-20.2500	-1.3737	1.8573	29.4830	27.0246
	GPT3.5-Turbo	-7.1533	-0.5109	1.7641	28.0048	20.4278
	llama2-70B	-3.8257	-1.4323	**0.3365**	**5.3420**	**4.1451**
	llama3-70B	-25.5829	-1.5579	2.0690	32.8437	31.3391
	gemini	8.6692	0.8015	1.3627	21.6321	18.4815
GM	Buy & Hold	0.3393	0.0179	2.3823	37.8181	23.0317
	GPT-4	10.5648	0.7671	1.7351	27.5443	11.1285
	GPT-4o	-7.0147	-0.5263	1.6792	26.6569	21.5978
	GPT3.5-Turbo	-17.6385	-0.9692	2.2928	36.3976	23.0317
	llama2-70B	8.4911	**2.6369**	**0.4057**	**6.4402**	**2.1318**
	llama3-70B	25.9335	1.9823	1.6483	26.1657	13.2485
	gemini	18.6257	2.4672	0.9511	15.0989	3.0369
NIO	Buy & Hold	-49.4263	-1.1895	5.2351	83.1048	52.2083
	GPT-4	**24.7684**	**0.9438**	3.3063	52.4861	29.3384
	GPT-4o	-48.3748	-1.5026	4.0562	64.3897	59.4037
	GPT3.5-Turbo	-28.9321	-1.0096	3.6105	57.3149	39.5907
	llama2-70B	-49.6947	-2.7868	**2.2466**	**35.6639**	42.6221
	llama3-70B	-28.6668	-0.7094	5.0912	80.8202	37.1544
	gemini	14.5673	0.6212	2.9543	46.8977	**23.0110**
COIN	Buy & Hold	-18.4787	-0.3369	6.9098	109.6904	60.5084
	GPT-4	25.7631	0.5619	5.7761	91.6934	35.7526
	GPT-4o	-14.2451	-0.2892	6.2049	98.4997	65.3090
	GPT3.5-Turbo	25.1141	0.4772	6.6312	105.2669	53.9628
	llama2-70B	15.1836	0.4395	**4.3528**	**69.0979**	**35.3249**
	llama3-70B	19.8876	0.3749	6.6842	106.1076	55.7225
	gemini	**89.4782**	**1.7648**	6.3879	101.4048	40.3246

F FinLLM challenge

Based on our proposed FinBen, we organized the FinLLM Share Task during the FinNLP-AgentScen Workshop at IJCAI 2024[9], known as the FinLLM Challenge. This challenge not only tests the

---

[9]https://sites.google.com/nlg.csie.ntu.edu.tw/finnlp-agentscen/shared-task-finllm?authuser=0

Figure 2: Accumulative Returns of LLM Trading Strategies on AAPL

Figure 3: Accumulative Returns of LLM Trading Strategies on AMZN

Figure 4: Accumulative Returns of LLM Trading Strategies on COIN


Figure 5: Accumulative Returns of LLM Trading Strategies on GOOG

abilities of LLMs but also promotes ongoing research into their application within the financial sector, highlighting FinBen's critical contribution to the advancement of financial analytics.

The FinLLM Challenge is a specialized shared task tailored for LLMs, targeting a comprehensive range of financial problems through three subtasks: financial classification, financial text summarization, and single stock trading. To rigorously evaluate the capabilities of financial LLMs, we have curated three distinct datasets corresponding to each of these subtasks, as detailed in Table 10. This
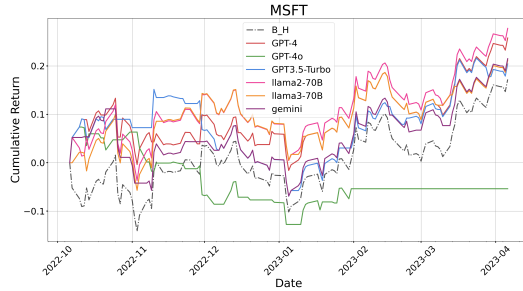
Figure 6: Accumulative Returns of LLM Trading Strategies on MSFT

Figure 7: Accumulative Returns of LLM Trading Strategies on NFLX

Figure 8: Accumulative Returns of LLM Trading Strategies on NIO


Figure 9: Accumulative Returns of LLM Trading Strategies on TSLA

structured approach ensures a holistic and effective assessment of LLM performance across diverse financial scenarios.

F.1 Tasks and Datasets

Task 1: Financial Classification. This task, inherited from FinBen's financial classification task, focuses on argument unit classification to test the capabilities of LLMs to identify and categorize
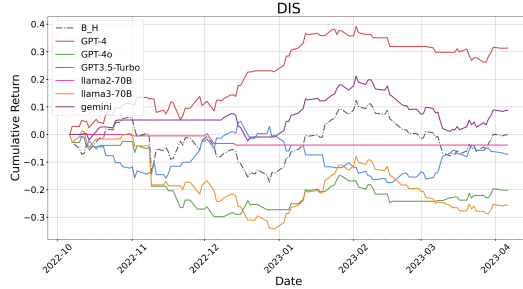
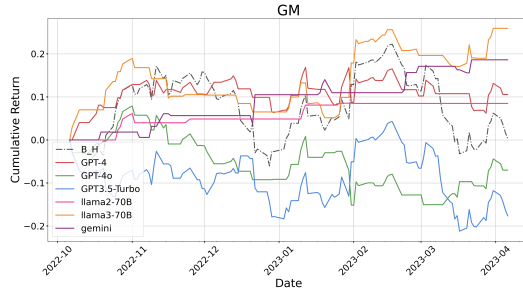Figure 10: Accumulative Returns of LLM Trading Strategies on DIS

Figure 11: Accumulative Returns of LLM Trading Strategies on GM

texts as premises or claims. It consists of 7.75k training data and 969 test data to categorize sentences as claims or premises. We use two metrics to evaluate classification capability, like F1 and Accuracy. F1 score is used as the final ranking metric.

Task 2: Financial Text Summarization. This task, inherited from FinBen's generation task, is designed to test the capabilities of LLMs to generate coherent summaries. It provides 8k training data and 2k test data for abstracting financial news articles into concise summaries. We utilize three metrics, such as ROUGE (1, 2, and L) and BERTScore, to evaluate generated summaries in terms of Relevance. ROUGE -1 score is used as the final ranking metric.

Task 3: Single Stock Trading. This task, inherited from FinBen's Trading task, aims to evaluate LLMs' ability to make sophisticated decisions in trading activities, which is currently restricted by human's limited ability to process large volumes of data rapidly. It specifically provides 291 data different from FinBen datasets, to evaluate LLMs on sophisticated stock Decisions. We offer a comprehensive assessment of profitability, risk management, and decision-making prowess by a series of metrics, such as Sharpe Ratio (SR), Cumulative Return (CR), Daily (DV) and Annualized volatility (AV), and Maximum Drawdown (MD). Sharpe Ratio (SR) score is used as the final ranking metric.

F.2 Model Cheating Detection

To measure the risk of data leakage from the test set used in training, we introduce the Data Leakage Test (DLT). The DLT calculates the difference in perplexity between the training set and the test set. A larger difference indicates a lower likelihood of model cheating, while a smaller difference suggests a higher likelihood. For our FinLLM Challenge, we invite Top-3 participant teams per task for cheating detection.

F.3 Participants and Automatic Evaluation

There are 35 teams registered for FinLLM Challenge, with 11 teams submitting their system description papers. Participants can opt to join one or more task(s).

As shown in Table 11, the top 3 teams achieved outstanding performance in Task 1. Their models' F1 scores were comparable to LlaMA3-8B, although slightly inferior to GPT-4 and LLaMA2-70B, yet

Table 10: Tasks and Datasets of FinLLM Challenge.

Category	Tasks	Datasets		Evaluation Metrics
		Training set	Test set	
Task 1	Financial Classification	7.75k	969	F1 Score, Acc
Task 2	Financial Text Summarization	8k	2k	ROUGE-1, ROUGE-2, ROUGE-L, BERTScore
Task 3	Single Stock Trading	291	225	Sharpe Ratio, Cumulative Return, Maximum Drawdown, Daily and Annualized Volatility,

significantly outperformed FinMA and other models. The results in Table 11 further demonstrate that our FinLLM share task provides an excellent framework for participating teams to achieve superior experimental outcomes.

Table 11: The Result of Taks 1: Financial Classification

Teams	ACC	F1	MCC
Team Barclays	0.7626	0.5237	0.7427
Albatross	0.7574	0.5174	0.7555
L3iTC	0.7544	0.5149	0.7581
Wealth Guide	0.7513	0.5018	0.7406
Finance Wizard	0.7286	0.4554	0.7008
CatMemo	0.711	0.4199	0.6818
Upaya	0.709	0.4166	0.6941
Vidra	0.7079	0.4141	0.69
jt	0.4933	0.0141	0.5905

As illustrated in Table 12, in terms of the Rouge-1 metric, the models of these three teams surpassed all other models, demonstrating superior performance. The results in Table 2 indicate that, for financial generation tasks, our provided dataset and model framework help participating teams leverage their strengths and achieve better outcomes.

Table 12: The Result of Taks 2: Financial Text Summarization

Teams	Rouge-1	Rouge-2	Rouge-L	BertScore	BartScore
Wealth Guide	0.308893532	0.179468097	0.281924302	0.85959909	-4.961457408
Albatross	0.369077581	0.201058395	0.322684316	0.872049115	-3.933526929
LBZ	0.534616211	0.358105428	0.492179554	0.911732047	-3.407560172
L3iTC	0.366093426	0.187210467	0.304610677	0.875037043	-4.257126737
Finance Wizard	0.521037018	0.34060938	0.473530112	0.90836845	-3.497988865
Vidra	0.284955468	0.134760859	0.228638961	0.858682767	-4.169740305
Revelata	0.500411369	0.333023818	0.464356474	0.907018743	-3.805486962
Upaya	0.529459817	0.358203218	0.486046685	0.910644962	-3.45155009

As shown in Table13, the Top-1 Wealth Guide team excelled in the Sharpe Ratio metric, surpassing other teams and demonstrating outstanding performance. While it may not match the performance of GPT-4, it still outperforms other large models. These results from Table 3 once again underscore the significance of organizing the FinLLM share task. The FinLLM Challenge not only assesses the performance of large language models (LLMs) but also fosters further research into applying LLMs in the financial domain.

G Performances of non-LLM methods

In this section, we present the performances of non-LLM methods on stock movement prediction and financial NLP tasks from previous papers. Note that non-LLM methods are task-oriented, each model can only run on a specific task.

G.1 Stock movement prediction

Stock movement prediction performance of non-LLM models are shown in Table 14. The results are from (Xie et al., 2023b).

G.2 Financial NLP tasks

BERT-based model results of financial NLP tasks are shown in Table 15. The results are from (Shah et al., 2022).

Table 13: The Result of Taks 3: Single Stock Trading

Teams	Sharpe Ratio	Sharpe Ratio-DRIV	Sharpe Ratio-FORM	Sharpe Ratio-JNJ	Sharpe Ratio-MSFT
Wealth Guide	0.9263852228	0.485625528	1.585611423	0.078737051	1.555566991
Upaya	0.467489019	0.380232272	0.108506918	-1.102831656	-0.278385232
Albatross	0.48383204	0.251306057	-1.435471054	-1.558522674	1.309971626
CatMemo	-0.619939784	-1.393291177	0.175932289	0.383243051	-0.879157198

Table 14: Stock movement prediction performance of non-LLM models, measured with the accuracy (ACC) and the Matthews correlation coefficient (MCC). The best performance is in bold.

Method	BIGDATA22		ACL18		CIKM18	
	ACC	MCC	ACC	MCC	ACC	MCC
Logistic regression (LR)	0.53	0.02	0.52	0.04	0.53	-0.04
Random forest (RF)	0.47	-0.11	0.52	0.03	0.54	0.01
LSTM	0.51	0.01	0.53	0.06	0.53	0.02
Attention LSTM (ALSTM)	0.49	-0.03	0.52	0.04	0.53	-0.01
Adv-ALSTM	0.50	0.01	0.53	0.07	0.54	0.02
DTML	0.52	0.07	0.58	0.18	0.54	-0.00
XGBoost	0.52	-0.04	0.49	-0.02	**0.58**	0.07
XGBRegressor	0.46	-0.13	0.50	-0.01	0.53	-0.03
ALSTM-W	0.48	-0.01	0.53	0.08	0.54	0.03
ALSTM-D	0.49	0.01	0.53	0.07	0.50	-0.04
StockNet	0.53	-0.00	0.54	-0.03	0.52	-0.02
SLOT	**0.55**	**0.10**	**0.59**	**0.21**	0.56	**0.09**

Limitations

Despite the novel efforts to benchmark LLMs in the financial domain through FinBen, we acknowledge several inherent limitations that could impact the benchmark's effectiveness and applicability:

Despite the novel efforts to benchmark LLMs in the financial domain through FinBen, several limitations could impact its effectiveness and applicability: **Dataset Size Limitations**: The restricted size of available datasets, a common issue in open-source financial data, may affect the models' financial understanding and generalization across various contexts. **Model Size Limitations**: Due to computational constraints, our evaluation was limited to the LLaMA 70B model, potentially overlooking the capabilities of larger or differently architected models. **Generalizability**: The tasks, particularly trading and forecasting, are based on American market data and English texts, possibly limiting the benchmark's applicability to global financial markets. **Potential Negative Impacts**: While FinBen aims to advance financial language understanding, it is crucial to consider potential misuse, such as propagating financial misinformation or exerting unethical influence on markets. Responsible usage and further safeguards are essential[10].

Ethical Statement

The development and dissemination of the FinBen by the authors carry full responsibility for any potential violation of rights or arising legal issues. All raw data we used are publicly available and do not contain any personal information. Diligent efforts have been undertaken to ensure the construction of the FinBen respects privacy and conforms to established ethical guidelines. The datasets compiled within FinBen are shared under the MIT license, with the expectation that users agree to adhere to its conditions.

This manuscript, inclusive of any associated source codes, datasets, and appendices ("Material"), is designated exclusively for academic and educational pursuits. It is crucial to acknowledge that the Material does not provide financial, legal, or investment counsel, nor should it be utilized as a foundation for any form of decision-making.

While the authors have exerted reasonable diligence to verify the accuracy and reliability of the Material, no explicit or implied warranty is extended regarding its completeness or suitability for any specific application. The authors, along with their affiliated entities, absolve themselves of liability for any losses, damages, or other consequences, whether direct or indirect, that may emanate from

---

[10]For a detailed ethical and legal statement concerning this work, please see Appendix.

Table 15: Financial NLP tasks performances of BERT-based models. The best performance is in bold.

Method	FPB Accuracy	Headline AvgF1	NER F1	FiQA SA MSE
BERT-base	0.856	0.967	0.79	0.073
FinBERT	0.872	0.968	0.8	0.070
FLANG-BERT	0.912	0.972	0.83	0.054
ELECTRA	0.881	0.966	0.78	0.066
FLANG-ELECTRA	0.919	0.98	0.82	0.034

the employment or reliance upon the Material. It is incumbent upon the user to seek professional consultation for financial, legal, or investment determinations.

By referencing or employing this Material, individuals consent to indemnify, defend, and hold the authors, along with any affiliated organizations or persons, harmless against any claims or damages that may arise from such utilization.